



Høyskolen
Kristiania

TK1100

Digital Technology

Computer architecture



Learning objectives covered so far (1)

- What is in the computer system
 - computer assembly and main components
- The computer language, alphabet,
- The Binary Numbering system.

Today: Computer Architecture

Boolean algebra

- NOT, AND, OR, XOR
- Sequential and combinatorial circuits.

Instruction Set Architecture (ISA).

- What instructions?
- Word size (number of bits)
- Addressing mode
- Register
- Data formats
- What the CPU is suitable of doing

Organization (microarchitecture on CPU).

- How data is being moved (buses architecture)
- How data is stored, e.g. Cache.
- ...

Now it is time to learn about computer architecture.
How the computer brain, CPU, in the hardware works with the alphabet.
how the spirit of numbers goes to the body of the computer architecture

Boolean algebra

Logic Gates

		y	
	\wedge	0	1
x	0	0	0
	1	0	1

		y	
	\vee	0	1
x	0	0	1
	1	1	1

		y	
	\rightarrow	0	1
x	0	1	1
	1	0	1

		y	
	\oplus	0	1
x	0	0	1
	1	1	0

Figure 1. Truth tables



Figure 2. Logic gates



Figure 3. De Morgan equivalents




Figure 4. Venn diagrams

A logic gate is an idealized model of computation or **physical electronic device implementing a Boolean function**, a logical operation performed on one or more binary inputs that produces a single binary output.

1	10
2	11
3	12
4	13
5	14
6	15
7	16
8	17
9	18

Bit and George Boole

- All logic in **digital** electronics is based on **bits**
- transistors are the building blocks
- Can be 0 or 1; high or low voltage.
- The logic used is called **Boolean Algebra**
 - George Boole (1815-1864)




Note! These 12 bits are normally just a snippet of a 32/64 bit word.



Boolean algebra


- Consists of 3 (4) basic operations (gates).
 - (NOT)
 - (AND)
 - (OR)
 - (XOR)
- Rediscovered in 1939 by Claude Elwood Shannon
- Computer electronics most often prefer «negative» logic
 - NAND, NOR, XNOR



he tried to solve the maze in one of the first experiments in artificial intelligence.

NOT


10010110



01101001

NOT
Truth table

X	F
0	1
1	0




- Flip-flops and latches are fundamental building blocks of digital electronics systems.
- Flip-flops and latches are used as data storage elements.
- A flip-flop is a device which stores a single bit (binary digit) of data

AND and NAND

10010101



10101100



10000100

AND


X	Y	F
0	0	0
0	1	0
1	0	0
1	1	1




NAND

X	Y	F
0	0	1
0	1	1
1	0	1
1	1	0

Example: AND




A	B	C
0	0	0
0	1	0
1	0	0
1	1	1




OR and NOR

10111001

10001110



10111111




OR

X	Y	F
0	0	0
0	1	1
1	0	1
1	1	1


NOR

X	Y	F
0	0	1
0	1	0
1	0	0
1	1	0



Example: OR



A	B	C
0	0	0
0	1	1
1	0	1
1	1	1



XOR and XNOR




XOR		
X	Y	F
0	0	0
0	1	1
1	0	1
1	1	0

XNOR		
X	Y	F
0	0	1
0	1	0
1	0	0
1	1	1

Example: XOR

A	B	C
0	0	0
0	1	1
1	0	1
1	1	0



Examples of Boolean algebra using binary system computer arithmetics

Binary Operations

- We learn the binary operations using truth tables

X	Y	AND
0	0	0
0	1	0
1	0	0
1	1	1

X	Y	OR
0	0	0
0	1	1
1	0	1
1	1	1

X	NOT
0	1
1	0

X	Y	XOR
0	0	0
0	1	1
1	0	1
1	1	0

- Given two bits, apply the operator
 - 1 AND 0 = 0
 - 1 OR 0 = 1
 - 1 XOR 0 = 1
- Apply the binary (Boolean) operators *bitwise* (in columns) to binary numbers as in
 - 10010011 AND 00001111 = 00000011

Examples

- AND – if both bits are 1 the result is 1, otherwise 0
 - $11111101 \text{ AND } 00001111 = 00001101$
 - $01010101 \text{ AND } 10101010 = 00000000$
 - $00001111 \text{ AND } 00110011 = 00000011$
- OR – if either bit is 1 the result is 1, otherwise 0
 - $10101010 \text{ OR } 11100011 = 11101011$
 - $01010101 \text{ OR } 10101010 = 11111111$
 - $00001111 \text{ OR } 00110011 = 00111111$
- NOT – flip (negate) each bit
 - $\text{NOT } 10101011 = 01010100$
 - $\text{NOT } 00001111 = 11110000$
- XOR – if the bits differ the result is 1, otherwise 0
 - $10111100 \text{ XOR } 11110101 = 01001001$
 - $11110000 \text{ XOR } 00010001 = 11100001$
 - $01010101 \text{ XOR } 01011110 = 00001011$

Binary Addition - (to remember)

- To add 2 bits, there are four possibilities
 - $0 + 0 = 0$
 - $1 + 0 = 1$
 - $0 + 1 = 1$
 - $1 + 1 = 2$ – we can't write 2 in binary, but 2 is 10 in binary, so write a 0 and carry a 1
- To compute anything useful (more than 2 single bits), we need to add binary numbers
- This requires that we chain together carries
 - The carry out of one column goes to be added in the other column to its left

Binary Addition Continued

- With 3 bits (the two bits plus the carry), we have 4 possibilities:
 - $0 + 0 + 0 = 0$
 - 2 zeroes and **1 one** = 1
 - **2 ones** and 1 zero = 2 (carry of 1, sum of 0)
 - **3 ones** = 3 (carry of 1 and sum of 1)
- Example:

Carry:		1	←	1	←	0	←	0	Initial carry
X:		0		1		1		1	in is 0
Y:	+	0		1		1		0	
Sum:		1		1		0		1	
				Carry		Carry		No Carry	


Check your work, convert to decimal!

Addition Using AND, OR, XOR

- To implement addition in the computer, convert addition to AND, OR, NOT and XOR
- Input for any single addition is two binary numbers and the carry in from the previous (to the right) column
 - For column i , we will call these X_i , Y_i and C_i
- Compute sum and carry out for column i (S_i , C_{i+1})
- $S_i = (X_i \text{ XOR } Y_i) \text{ XOR } C_i$
 - Example: if $1 + 0$ and carry in of 1
 - $\text{sum} = (1 \text{ XOR } 0) \text{ XOR } 1 = 1 \text{ XOR } 1 = 0$
- $C_{i+1} = (X_i \text{ AND } Y_i) \text{ OR } (X_i \text{ AND } C_i) \text{ OR } (Y_i \text{ AND } C_i)$
 - Example: if $1 + 0$ and carry in of 1
 - $\text{carry out} = (1 \text{ AND } 0) \text{ OR } (1 \text{ AND } 1) \text{ OR } (0 \text{ AND } 1) = 0 \text{ OR } 1 \text{ OR } 0 = 1$
- Try it out on the previous problem

Practical example: Half adder


- Adds two bits, A and B
- Gets the answer in **Q** and carry in **CO**
 - $Q = A \text{ XOR } B$
 - $CO = A \text{ AND } B$



A	B	CO	Q
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0

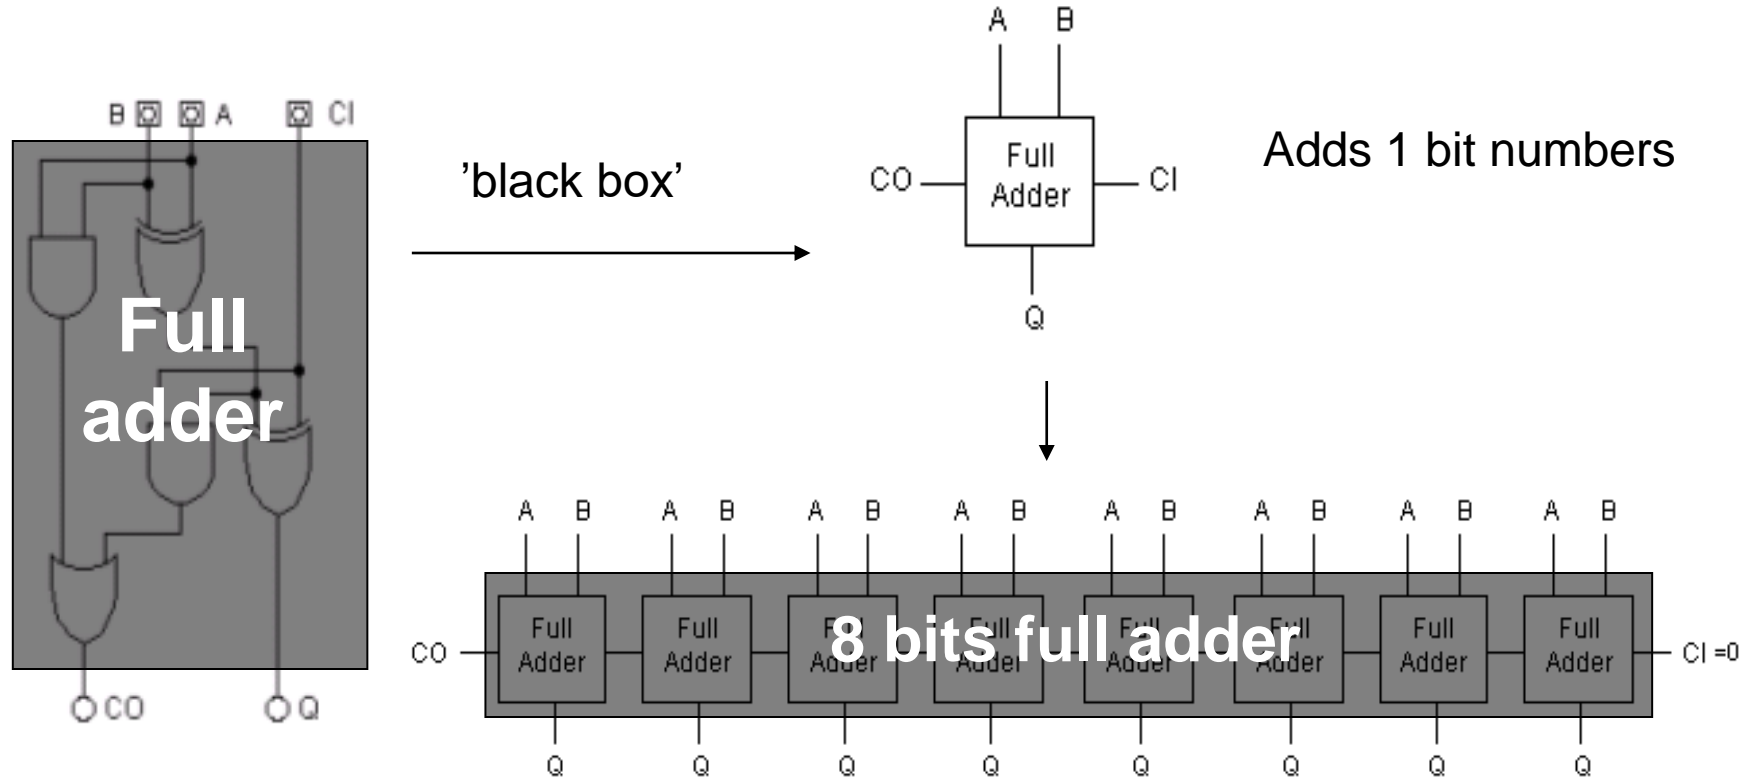
Practical example: Full adder

- Adds two bits together, **A** and **B**, and possibly **CI** (carry in)
- Gets the answer in **Q** with carry in **CO** (carry out)



A	B	CI	Q	CO
0	0	0	0	0
0	0	1	1	0
0	1	0	1	0
0	1	1	0	1
1	0	0	1	0
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1

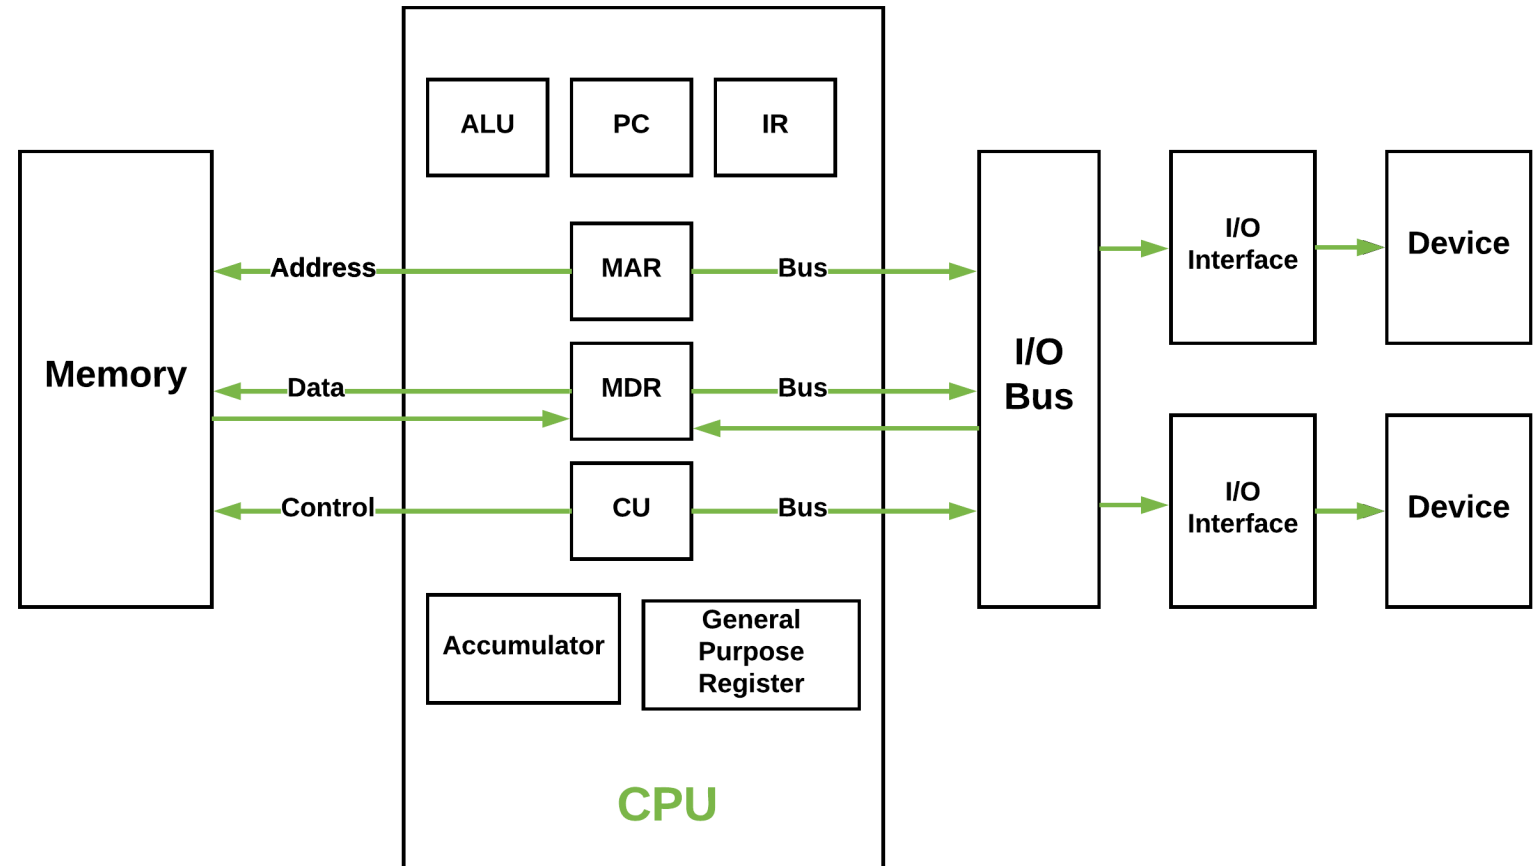

Practical example: 8 bit full adder



8 full adders connected together: Adds 8 bit numbers

Computer architecture

Von Neumann model



<https://youtu.be/-SADbPS8UgA>

C_{entral} P_{rocessing} U_{nit}

- Areas of use
- Working speed (clock frequency)
- Instruction set (ISA)
- Register
- Arithmetic Logical Unit and Float Point Unit
- Techniques to increase performance: pipelining and parallelization

Processor Classification

- Instruction **S**et **A**rchitecture (**ISA**)
 - RISC vs CISC
- Area of use
 - Embedded <-> Desktop <-> supercomputers
- Clock frequency
- **Registry file size**
- Number of functional units **arithmetic logic unit (ALU)** and floating-point unit (**FPU**)
- Superscalarity and parallelism
- Socket-type
- Energy consumption and cooling needs

Different Processor (CPU) Design Methodologies

- RISC vs CISC
- **C**omplex **I**nstruction **S**et **C**omputing
- **R**educed **I**nstruction **S**et **C**omputing



Complex Instruction Set

0 0 0 1 0 1 0 1 0 0 1 0 1 1 0 0 0 0 0 0 1 1 1 0 0 0 1 0 1 0 0 1

Operation code

Operand

Reduced Instruction Set

0 0 1 1 1 0 1 1 0 1 1 0 0 0 0 1 0 1 1 0 0 0 1 1 0 1 0 0 0 1 1 0

Operation code


Operand

<https://www.youtube.com/watch?v=6Rxade2nEjk&t=8s>

Binary number representation (recall)

- With 3 lamps we can represent $2^3 = 8$ combinations of on / off
- With Off = 0 and On = 1 we get

0 = 000	4 = 100
1 = 001	5 = 101
2 = 010	6 = 110
3 = 011	7 = 111
- This can be extended to use more lamps (transistors), e.g. 8, 16, 32, 64,



Manufacturers-Instruction Set Architecture (ISA)

Mac book
pro intel

- Intel (x86) and AMD

- **C**omplex **I**nstruction **S**et **C**omputing
 - Constructs and produces CPU, chipset etc.
 - Focus: **Desktop**, **server**, mobile, embedded.

other Mac
apple M1
ARM

- MIPS(32 registers) and ARM (16 registers)(apple)

- **R**educed **I**nstruction **S**et **C**omputing
 - Constructs and licenses CPU cores etc. to various manufacturers
 - Focus: **Embedded**, **mobile**
- Now x86 and ARM processors both have features of RISC and CISC

Complex Instruction Set (CISC)



Reduced Instruction Set (RISC)



<https://www.youtube.com/watch?v=6Rxade2nEjk&t=8s>

Area of use

- Embedded systems
- IoT Devices
- Supercomputer
- Mobile phones use simple and "cheap" CPUs.
- Desktop and Laptops

Embedded and IoT systems


IoT vs Embedded Systems

	IoT	Embedded Systems
Definition	Things that contain computers, software and networking capabilities.	Things that contain computers and software.

- Made for specific purposes
- Controlled by a microprocessor /- **controls**
- The software is ready-made, often a single program, and is usually in firmware.
 - Can be configured frequently.
 - Does not (often) need OS.
- Example
 - Washing machines, home routers, clock radios, guitar tuners, smart cards, SIM cards,...

- In the last couple of years, the **Arduino** and Raspberry Pi have opened up for much simpler home development of various embedded systems.

<https://www.youtube.com/watch?v=p40OetppIDg>



Supercomputers



Kilde: <http://top500.org/>

Tianhe-2

- 3.120.000 cores
- 33.863 Tflop/s
- 17.808 kW
- 1.024.000 GB RAM
- Kylin Linux

The world's
largest and
fastest: CHINA

Mobile phones

- Traditionally, mobile phones use simple and "cheap" CPUs.
 - **ARM** or **MIPS** ISA and core
 - The mobile manufacturers then produce custom versions themselves.
- Intel has ambitions in the market.
- "Smart phones" are (perhaps) most similar to a full-fledged PC with its own GSM card



Desktop \approx x86

- **x86**
 - x86 is a generic term for CPUs with the same set of instructions as Intel developed.
- **Apple** started with Motorola 68000 RISC CPU, continued with IBM / Freescale PowerPC CPUs.
 - Went over to x86 i 2005
- **Intel and AMD**
 - Still developing the x86 series.

Area of use

- Embedded systems
 - The dominant CPU architecture for embedded
 - ARM, which is always described as RISC.
 - Other players are x86 [CISC] and MIPS [RISC].
- IoT Devices
 - Most Arduinos (ARM, etc) are considered RISC
- Supercomputer
 - RISC offers some advantages over CISC and have even been used in the current world's fastest supercomputer (February 10, 2020).
- Mobile phones use simple and "cheap" CPUs.
 - ARM or MIPS ISA and core
- Desktop and Laptops


x86 instruction set (CISC)

- **Data relocation:**
 - Move data between registers and to / from RAM.
 - move, push, pop, ...
- **Control transfer:**
 - Perform condition statements (if, while) and method calls
 - je, jg, jmp, call, ret
- **Arithmetic / Logic:**
 - Perform operations on data in registers.
 - cmp, add, sub, inc, mul, imul, not, and, or, xor
- **Input/Output:** in, out
 - Writing / reading to ports (and memory addresses)
- **Debug and Interrupt handling:**
 - int, sti, hlt, nop
- **Floating point numbers** has their own instructions
- **Single instruction, multiple data (SIMD)** vector and matrix instructions.
 - MMX, 3D Now!, SSE 1-4
- Separate instructions for changing the processor state (system)




Techniques and concepts

System clocks (Hz)

- The system clock keeps steps, does not show the time!

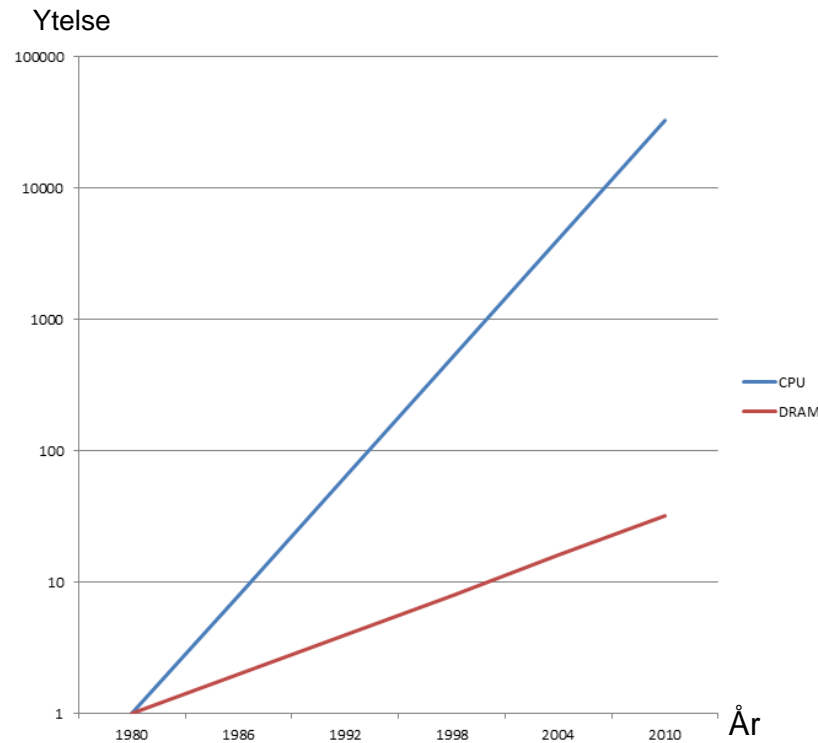


3.2 GHz executes 3.2 billion cycles per second. (Older CPUs had speeds measured in megahertz, or millions of cycles per second.)



CPU vs RAM (DRAM)

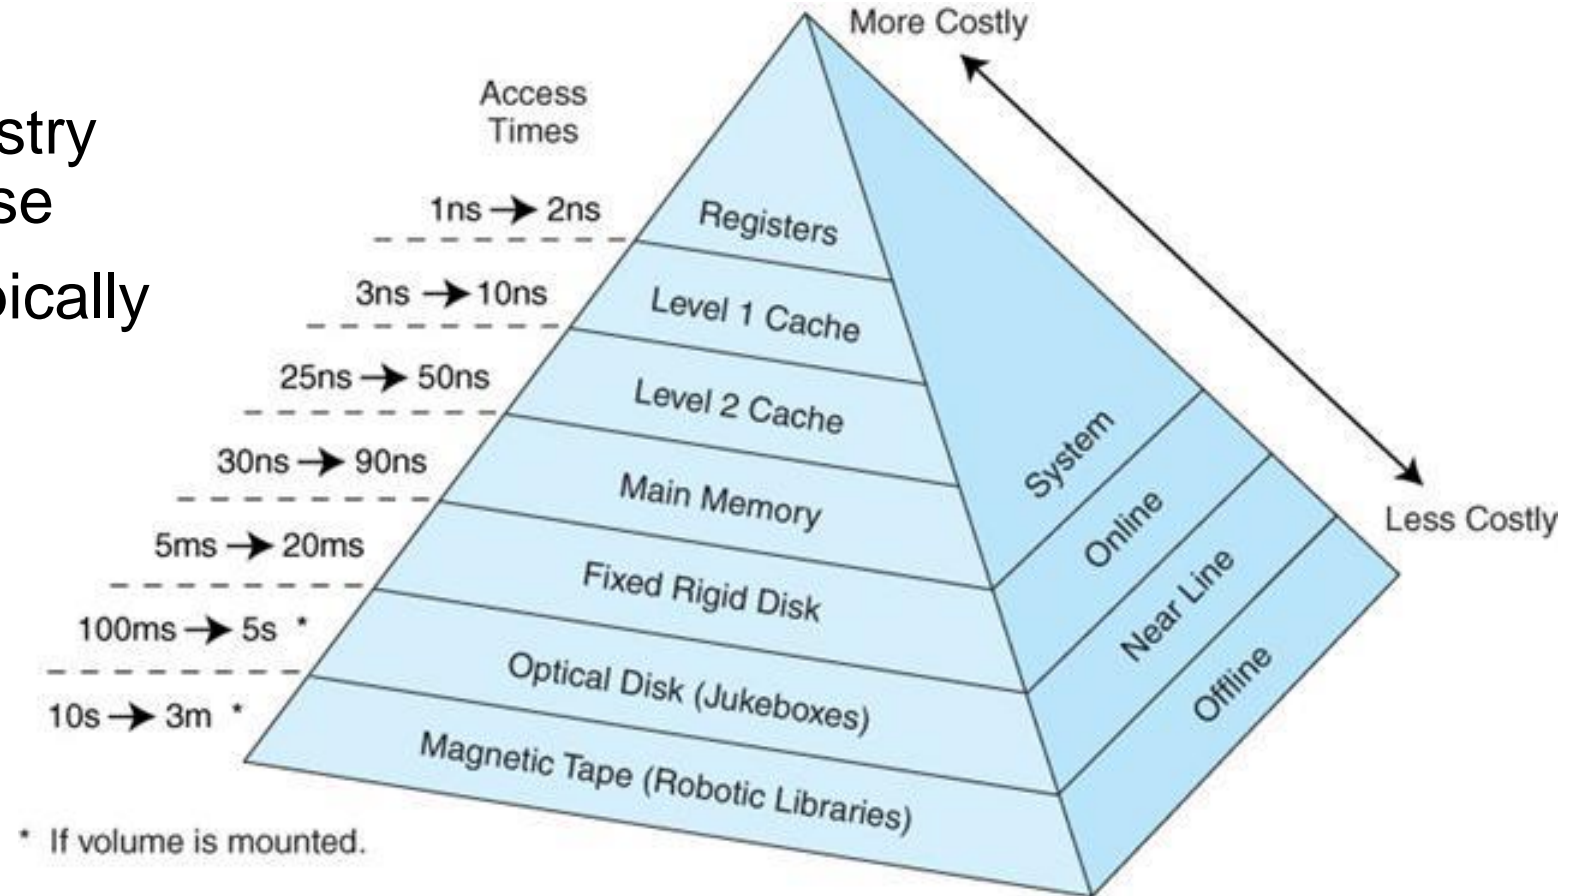
- CPU performance / speed have doubled every other year
- DRAM performance / speed every six years!



- The complete program is in RAM
- Only single instructions and single data are on the CPU
- -> PROBLEM !!!
- The challenge is to keep the pipeline full

Memory-hierarchy

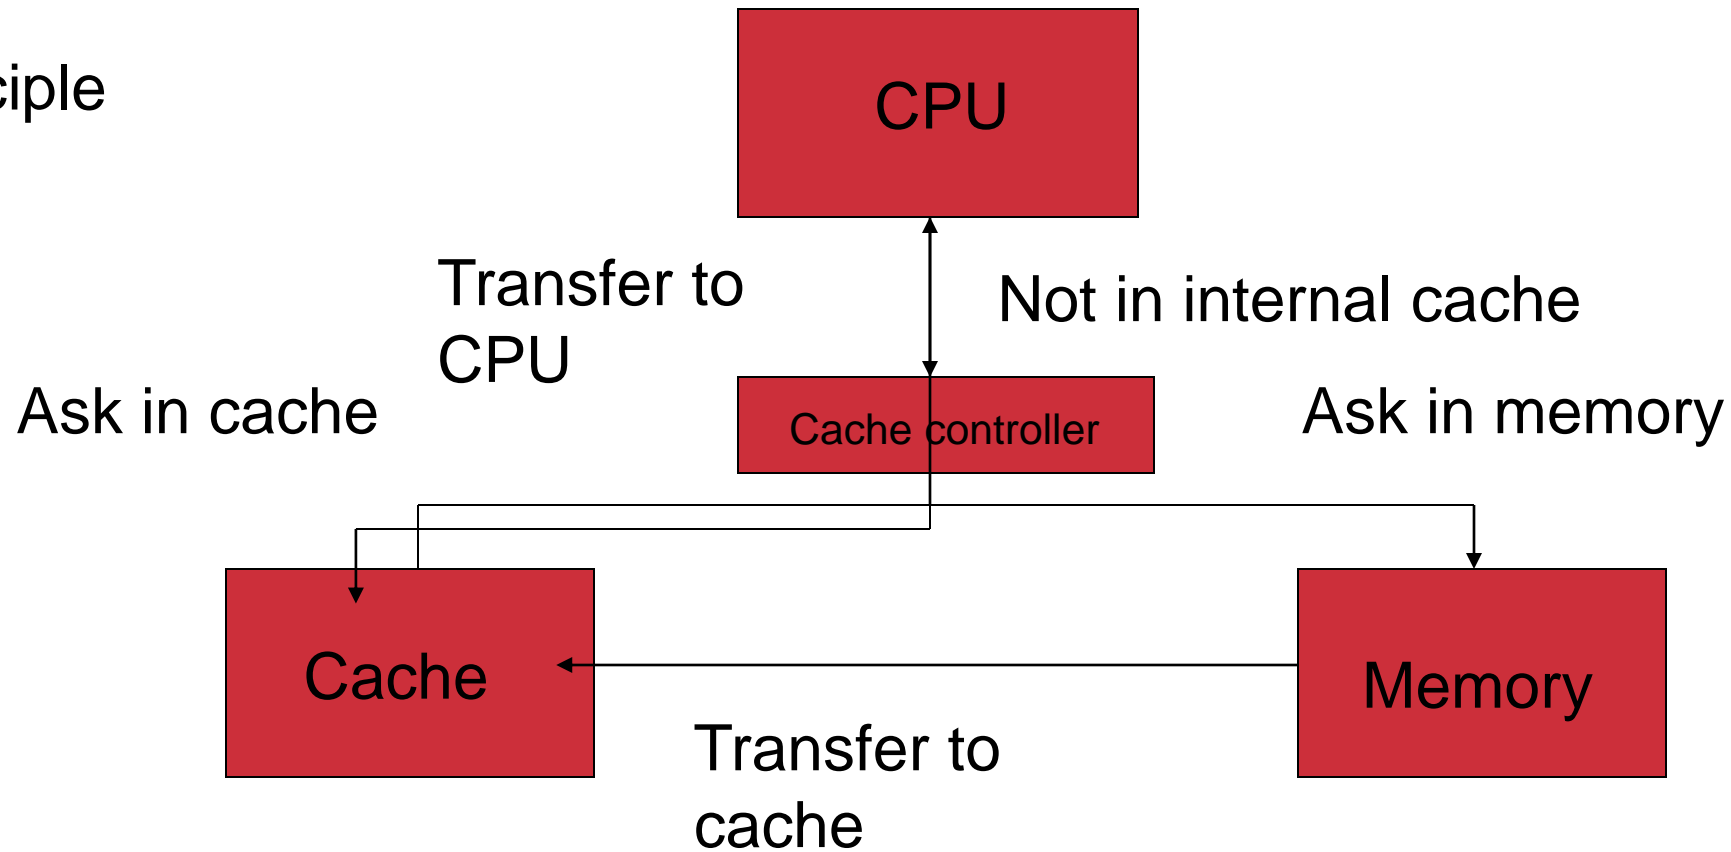
- Doing something in a registry typically takes a clock pulse
- Reading / writing RAM typically 150-300.



- **The locality principle**
 - About 5% of the code is run about 95% of the time.
 - Instructions to be executed are stored in the order in which they are going to be executed.
- Cacheing takes advantage of this by **caching** the most recently read **instructions**
 - **Examples:** keeping loops make it easy to access ...
- + Pageing
 - We typically do not read individual instructions to Cache, but memory blocks of 4 KiB

Cache

- In principle



More "cores"

- A **multi-core processor** is a computer processor on a single integrated circuit with two or more separate processing units, called cores, each of which reads and executes program instructions
- The instructions are ordinary CPU instructions (such as add, move data, and branch) but the single processor can run instructions on separate cores at the same time, increasing overall speed for programs that support multithreading or other parallel computing techniques.
- The CPU contains several full-fledged "sub-CPU's", which can run software in parallel.
- you must have programs that can be parallelized and must not be run sequentially!




Diagram of a generic dual-core processor with CPU-local level-1 caches and a shared, on-die level-2 cache

Units and sizes

Bit (b) - 0 or 1

Byte (B) = 8 bits

Kilo = $10^3 = 1000 \approx 1024 = 2^{10}$

1 km = 1000 m, 1 mm = 1/1000 m

For the sake of simplicity, we "cheat" a little!

Heart (Hz) = events per second

MIPS = Mega instructions per second.

FLOPS = Mega floating point operations per second

kbps = 1000 bit / second (bit rate, «bandwidth»)

System of Units (SI)			Binary Numeral				%
Factor	Name	Symbol	Factor	Name	Symbol	# of Bytes	Difference
10^3	kilobyte	KB	2^{10}	kibibyte	KiB	1,024	2.4%
10^6	megabyte	MB	2^{20}	mebibyte	MiB	1,048,576	4.9%
10^9	gigabyte	GB	2^{30}	gibibyte	GiB	1,073,741,824	7.4%
10^{12}	terabyte	TB	2^{40}	tebibyte	TiB	1,099,511,627,776	10.0%
10^{15}	petabyte	PB	2^{50}	pebibyte	PiB	1,125,899,906,842,624	12.6%
10^{18}	exabyte	EB	2^{60}	exbibyte	EiB	1,152,921,504,606,846,976	15.3%
10^{21}	zettabyte	ZB	2^{70}	zebibyte	ZiB	1,180,591,620,717,411,303,424	18.1%
10^{24}	yottabyte	YB	2^{80}	yobibyte	YiB	1,208,925,819,614,629,174,706,176	20.9%

Differences on x86 16, 32 and 64 bit?

- 16 bit architecture
 - 16 bit register size (word)
 - 20 bit addressing (max 1 MiB «RAM»)
 - Segmented addressing (CS: IP e.g)
 - 16 bit words (typically 16 bit int e.g)
 - No distinction between running kernel code in OS and regular programs.
- 32 bit (386->Pentium)
 - 32 bit register (word size)
 - 32 bit addressing (4 GiB(**gibibyte**) Virtual Memory)
 - Flat and support for virtual memory
 - 32 bit words
 - Distinguish between kernel and user runtime.
- 64 bit (now)
 - 64 bit registers and multiple work registers
 - Typical 48 bit addressing (16 EiB(**exbibyte**) virtual memory)
 - Often multi-core
 - Even more instructions

A mebibyte is equal to 2^{20} or 1,048,576 bytes. A megabyte is equal to 10^6 1,000,000 bytes. **One mebibyte equals 1.048576 megabytes.**

The relatively small difference between them and so they are often used synonymously.

End

Point

- CISC
 - COMPLEX Instruction Set Computer
 - Many different instructions
 - Instructions have different lengths in bits / bytes
- RISC
 - REDUCED Instruction Set Computer
 - Only those instructions that (empirically) are used a lot.
 - Provides simpler electronics and thus faster execution.
 - See byte code on JVM.
- Moderne Intel/AMD CPU
 - "both"
 - CISC externally
 - RISC microinstructions on the processor itself

More points

- The processor executes the instructions in the order in which they are stored in memory («**sequentially**»).
- On **different data types** completely **different instructions** are used
 - $7 + 5$:
`mov AX, 7`
`add AX, 5`
results in 12 being in the AX register
 - $7.0 + 5.0$:
`FLD «The memory address where 7.0 is located»`
`FLD 0012AC45`
`FADD ST, ST(1)`
`FSTP «the address the answer should be added»`
 - 'Hei' + 'Hå!'
loads of instructions that move both strings into the common memory area and add the end cursor (typically ASCII 0x00).

More points

- What is a variable??
 - At the instruction level: **the address** of a place memory (or a register).
- Why do variables have «scope»?
 - Calling a method means that you start running instructions somewhere else in memory
 - When you jump back to where the method was called, everything other than what the method returns disappears (typically an address to a value...)

What should we know?

- Boolean algebra
 - NOT, AND, OR, XOR
- What an instruction is.
 - how built up
 - what types
- What a register is.
 - the types of registers found on the CPU
 - The difference between address and answer register.
- Von Neumann model (again)
 - smaller «black box»
- The difference between CISC and RISC
- A little about x86 and options

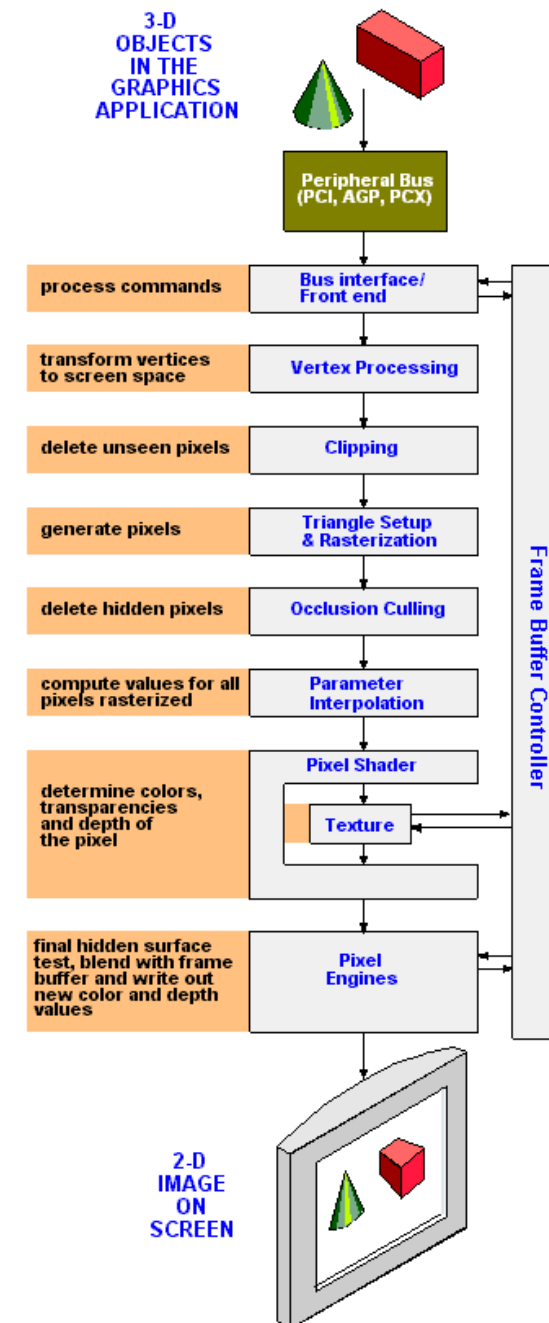
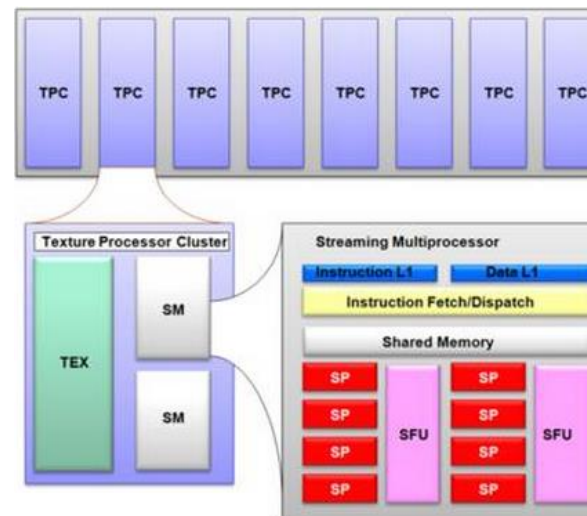
For optional self-study

For those who want to learn some topics more in depth to better understand, here are some extra topics related to today's teaching, it must be expected some personal work to understand these topics.

There will be no questions on the exam from these, and this is therefore not considered to be part of the syllabus.

Graphics Processing Unit


- Either integrated on motherboard, or mounted on graphics card
- Massively parallel
 - (now) ~ 100s of parallel arithmetic units
 - Calculates matrices (tables) of floating point numbers
 - Transforms mathematical models into screenshots (pixels in framebuffer)





Heavy rain

As of June 2014, there were various Intel Xeon models, which run Linux in Clusters dominate

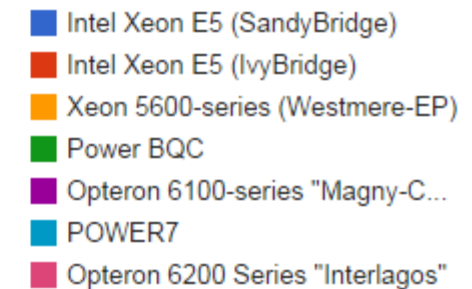

Application Area Performance Share



Operating system Family System Share



Processor Generation System Share




Software vs Hardware

Applications

- Shell + file manager
- User programs
- Compilers

Hardware


- CPU, RAM, I/O-controllers
- Peripheral buses
- Additional card
- Hard disk ...




Operating system

- Manage access to CPU, Memory, Equipment and files.
- Makes it easier to use hardware.
- Abstraction layer

Motherboard approx 2004 (?)

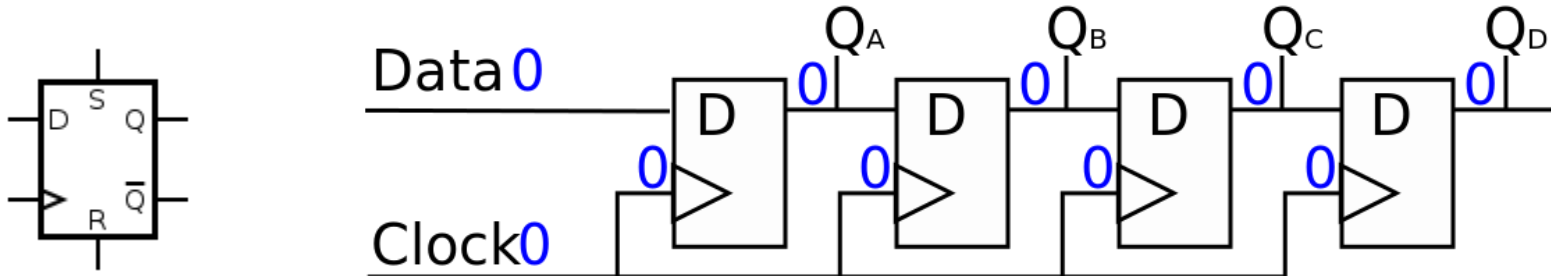


Motherboard 2012



Sequential vs. Combinatorial

- So far we have only looked at what is called **combinatorial** logic
- Output is determined solely by **Input** (and time).
- **Sequential** logic also has «memory»
 - Output is determined by **Input and stored bits**.
 - Used i.a. to build register (memory cells).



Boolean as data type

- In many contexts, there is no correspondence between the values 0 and 1 and Boolean data types.
 - Often used for statement logic, where we are interested in "true" and "false", not 1/0
 - The statement `4 == 4` has truth value **true**
 - The statement `3 == 5` has truth value **false**
 - How "true" and "false" are actually coded binary is **language / compiler** dependent.
- In Java (and other C-languages) you have different logical operators depending on whether it is truth value or bit-combination you want to perform.
 - **&** or **bitwise**, **&&** are **logical**
 - Thus becomes **`3&4 = 0`** (`0011 & 0100 = 0000`),
 - But **`3&&4`** gives «true», «syntax error», or 4...

Programming

Bitwise AND	writes	&
Bitwise OR	writes	
Bitwise XOR	writes	^
Bitwise NOT	writes	~
Logical NOT	writes	!
Logical AND	writes	&&
Logical OR	writes	

LET US DESIGN A SIMPLE COMPUTER!

- SIMPLE!!! (Understand the principles)
- Primitive
- We do not make electronics
- The machine can still do some things
- Further description can be found in **Compendium 1**

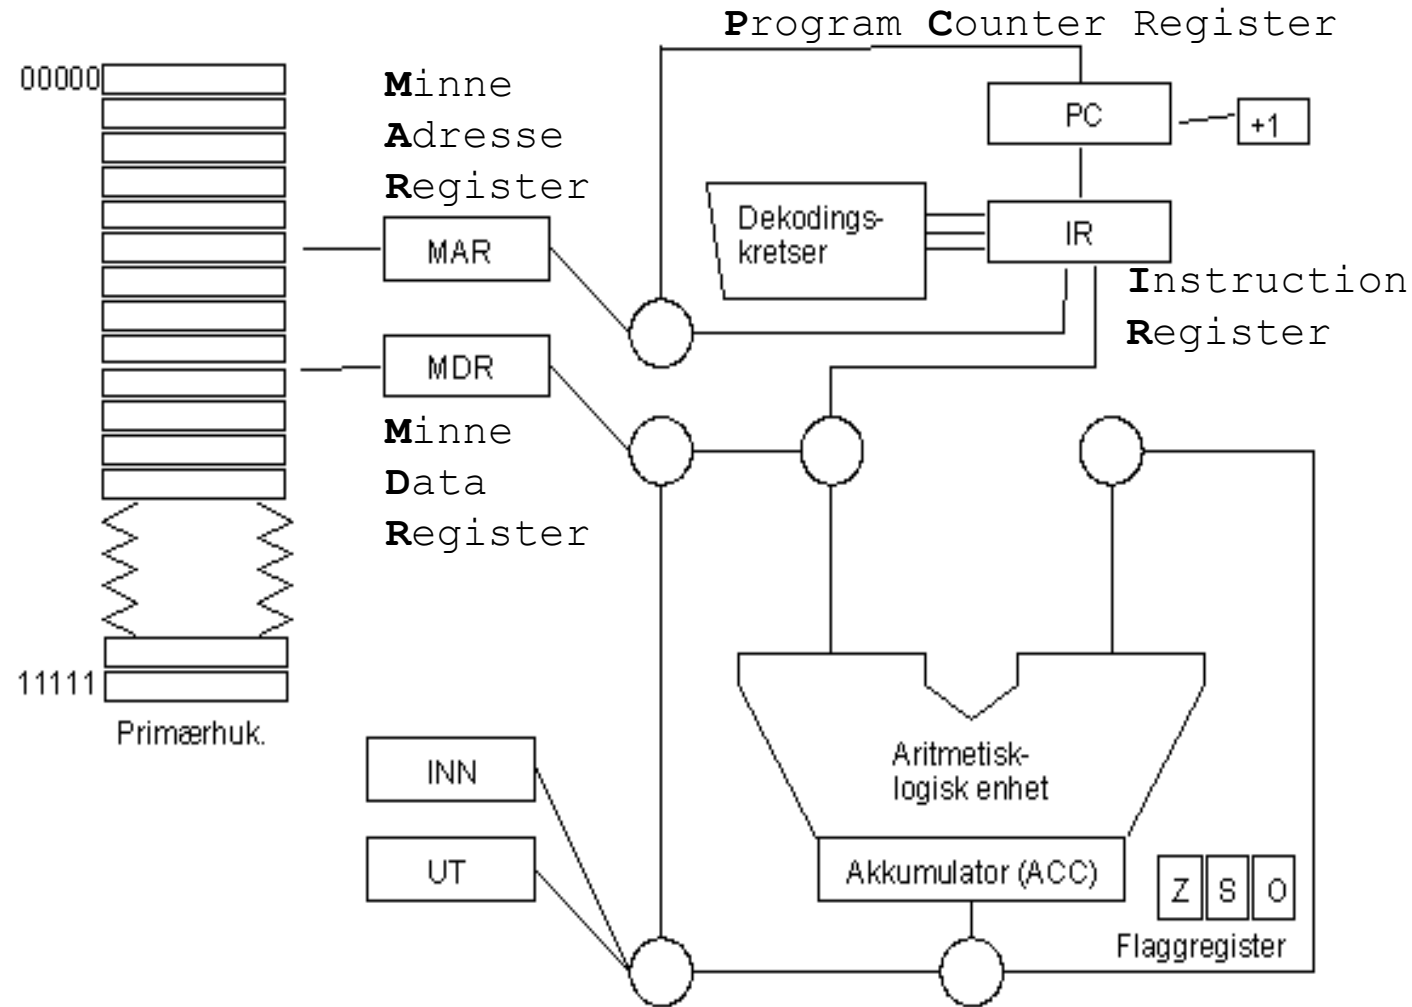
What was architecture again?

- How to build the instructions that the computer should execute (execute)
- What types of instructions should we have?
- How should addressing take place?
- Once we have determined the **I**nstruction **S**et **A**rchitecture, we can design a CPU
- In practice, this is an interaction between machine design and ISA design
- We have a transistor budget

Register

- All bits on the CPU must be entered in a register
- A register is addressable memory (SRAM) that resides on the CPU.
- In Assembly language, the registers typically have their own "names" and roles
- The registers used to manipulate variables and values are referred to as work registers.
- Example. In X86-64
 - **RIP** is the register that contains the address of where in memory the next instruction is to be retrieved from.
 - **RAX** is the register most often used to store results in (accumulator)
 - **RCX** most often used for countdown (think for-loop)
 - + many more (memory processing, driving,...)

PrimRisc v. 0.1 layout



- Accumulator-based
- Only one work register
- 8 bit addressing
- Uses only 5
- Load / store
- Everything goes via RAM
- Single
- Can "really" just add

Register- «file»

- **PC (Program Counter).**
 - This register will at all times contain the address of the location in the primary warehouse from which the next instruction is to be retrieved.
- **IR (Instruction Register).**
 - This register stores instructions after they have been retrieved from the primary memory and before being sent to the decoder circuits.
- **MAR (address register).**
 - When we want to read something from, or put something in the primary warehouse, we must first place the address in MAR.
- **MDR (data register).**
 - When we do a read operation from the primary memory, the result is placed here. If we are to write to the primary memory, we must first place what we are to write in MDR
- **ACC (Accumulator).**
 - This register is used primarily in connection with calculation operations. It is also possible to read the contents of a memory address into the accumulator with the instruction LOAD, and save the contents of the accumulator with the instruction STORE.
- **FLAG register**
 - we need it in order to be able to make conditional jumps (if, while and the like)
 - **The Z flag** (Zero flag) is set to 1 if the result of a calculation operation is zero.
 - **The S-flag** (Sign-flag) is set to 1 if the result of a calculation operation is negative.
 - **The O-flag** (Overflow-flag) is set to 1 if the result of a calculation operation becomes too large (overflow) to fit in the ACC register

Instructions

Mnemonic	Opkode	Beskrivelse	Funksjon
STOP	0000	Stopp maskinen	
LOAD	0001	Les inn i akkumulator	$(ACC) \leftarrow (MDR)$
STORE	0010	Lagre akkumulator	$(MDR) \leftarrow (ACC)$
ADD	0011	Adder til akkumulator	$(ACC) \leftarrow (ACC) + (MDR)$
SUB	0100	Subtraher fra akkumulator	$(ACC) \leftarrow (ACC) - (MDR)$
BRANCH	0101	Hopp uansett	$(PC) \leftarrow IR<4:0>$
BRZERO	0110	Hopp hvis resultat = 0	$(PC) \leftarrow IR<4:0>$ eller $(PC) \leftarrow (PC) + 1$
BRNEG	0111	Hopp hvis resultat ≤ 0	$(PC) \leftarrow IR<4:0>$ eller $(PC) \leftarrow (PC) + 1$
BROVFL	1000	Hopp hvis overflyt	$(PC) \leftarrow IR<4:0>$ eller $(PC) \leftarrow (PC) + 1$
IN	1001	Les inn fra tastatur	$(ACC) \leftarrow INN$
OUT	1010	Skriv til skjerm	$UT \leftarrow (ACC)$

- All instructions are 16 bit (2 bytes)
- 1 byte to **encode**
- Uses only 4 bits (in this version)
- 1 byte in the instruction to addresses
- Uses only 5 bits (in this version)

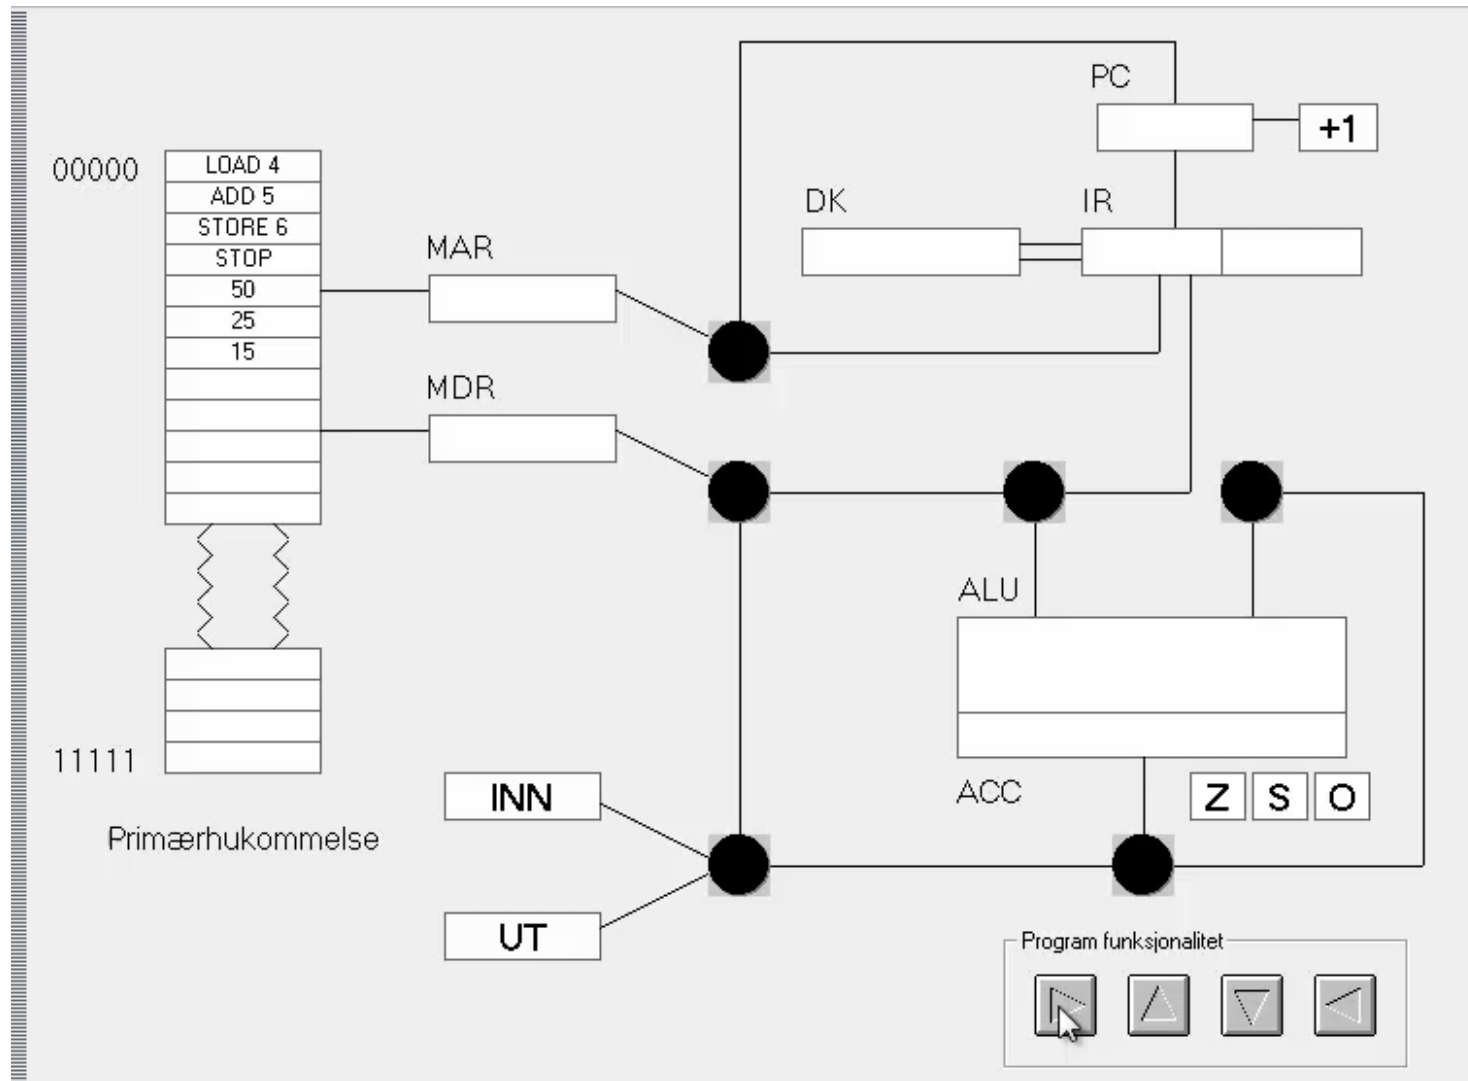
Example of a **program** in memory

Minneadresse binært/(desimalt)	Maskinkode (opkode+adresse) binært/(desimalt)	Asembly- instruksjon
00000 (0)	0001 0000 0000 0100 (1) (4)	LOAD 4
00001 (1)	0011 0000 0000 0101 (3) (5)	ADD 5
00010 (2)	0010 0000 0000 0110 (2) (6)	STORE 6
00011 (3)	0011 0000 0000 0000 (0) (-)	STOP
00100 (4)	0000 0000 0000 0001 (0) (1)	
00101 (5)	0000 0000 0000 0010 (-) (2)	
00110 (6)	0000 0000 0000 0000 (-) (0)	

```
/* Load in the accumulator data from address 4,  
add data from address 5, store at the address  
located at address 6,  
End the run */
```

- *Can use a simulator of this architecture (buggy) on today's exercise (Take a look at http://youtu.be/igthwn_qGVI)*

Test



Analysis / Criticism of PrimRISC 1


- Ready to demonstrate a **Fetch-Execute** cycle and how instructions can be structured
- Missing **Pipelining**
- Poor instruction set
- 8 against x86-64 sine ~ 1000
- Does not support System Programming (OS), multitasking or virtual memory
- «Quite small, but still so simple»

Criticism 2

- Has support for conditional jumps
- Can implement if, for, while etc.
- Missing (proper) HW support for methods !!!
- In Intel / AMD, this is solved with separate registers (SS and SP) and instructions (call) that support the use of a separate memory area as a stack.
- Does not support interrupt!
- More about this in the upcoming lectures.

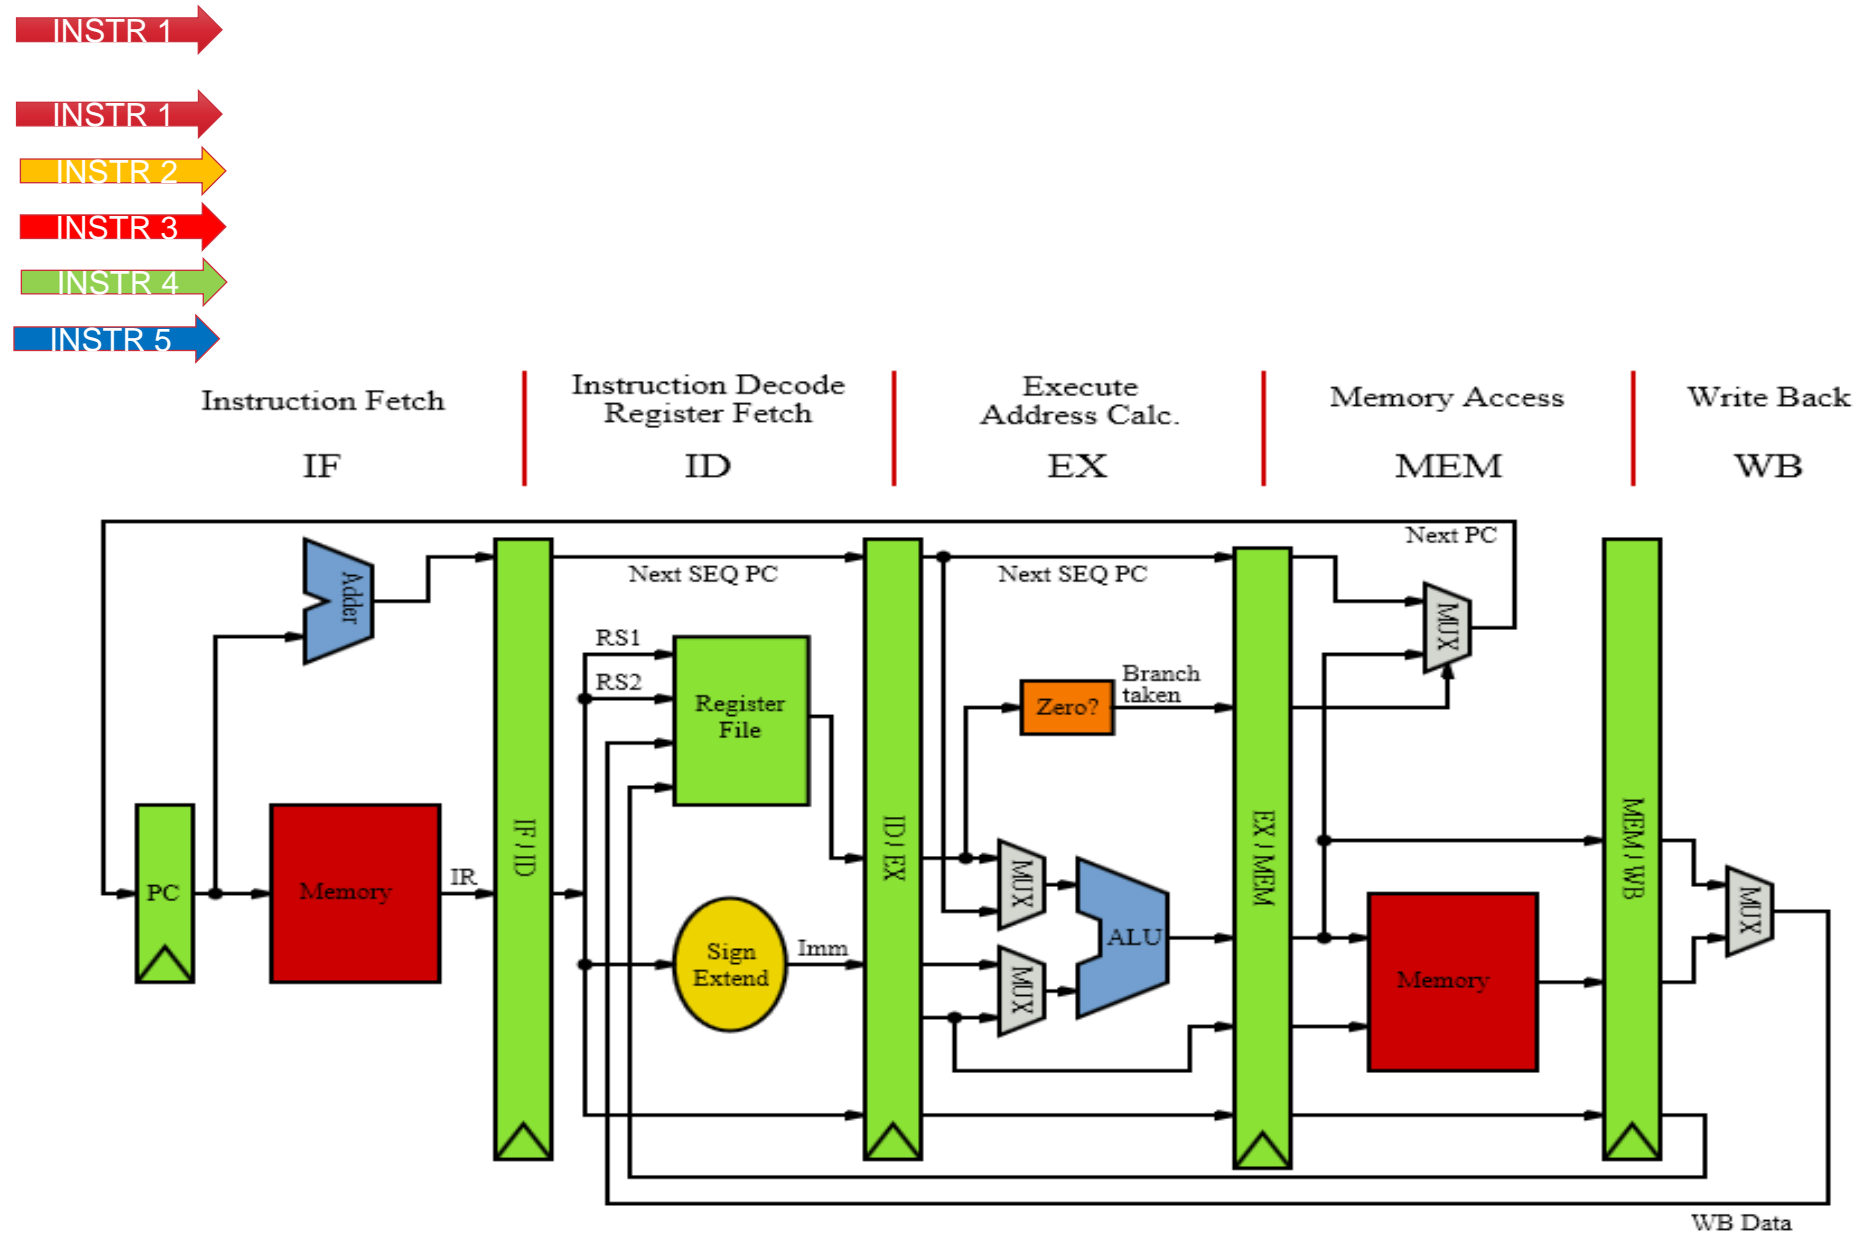
Fetch Execute Cycle

1. **Fetch:** retrieve instruction from memory and into registry
2. **Decode:** Interpret the instruction and set up the circuits that execute it
3. **Execute:** Execute the instruction on data / registers (change state)
4. **Store:** Save result (back in memory)

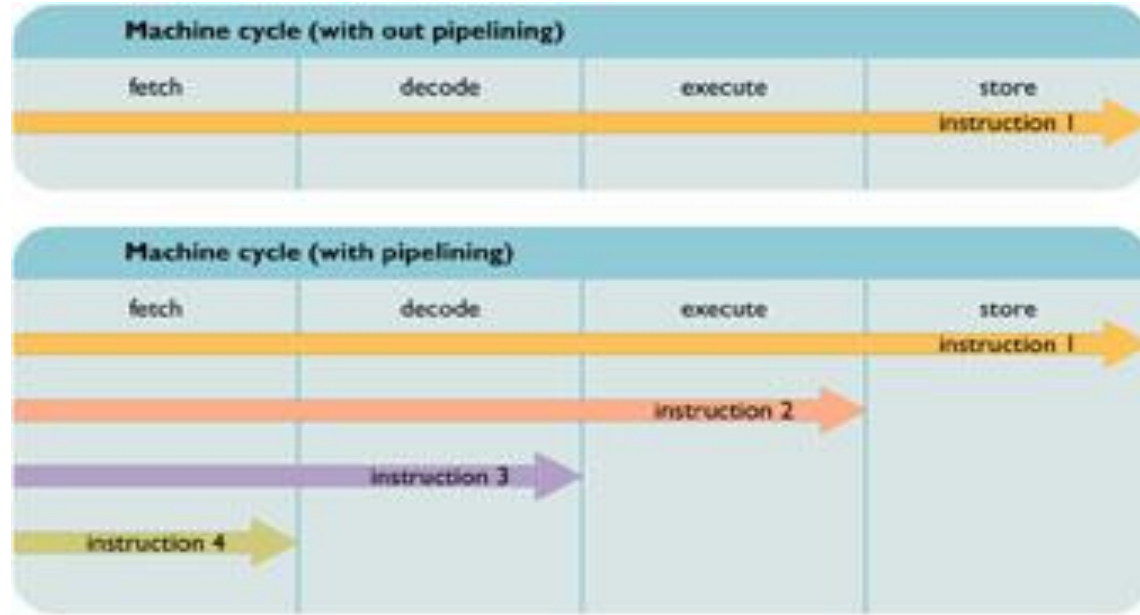



Instruction Fetch (IF)
Instruction Decode (ID)
Data Fetch (DF)
Instruction Execution (EX)
Result Return (RR)

Ex: MIPS

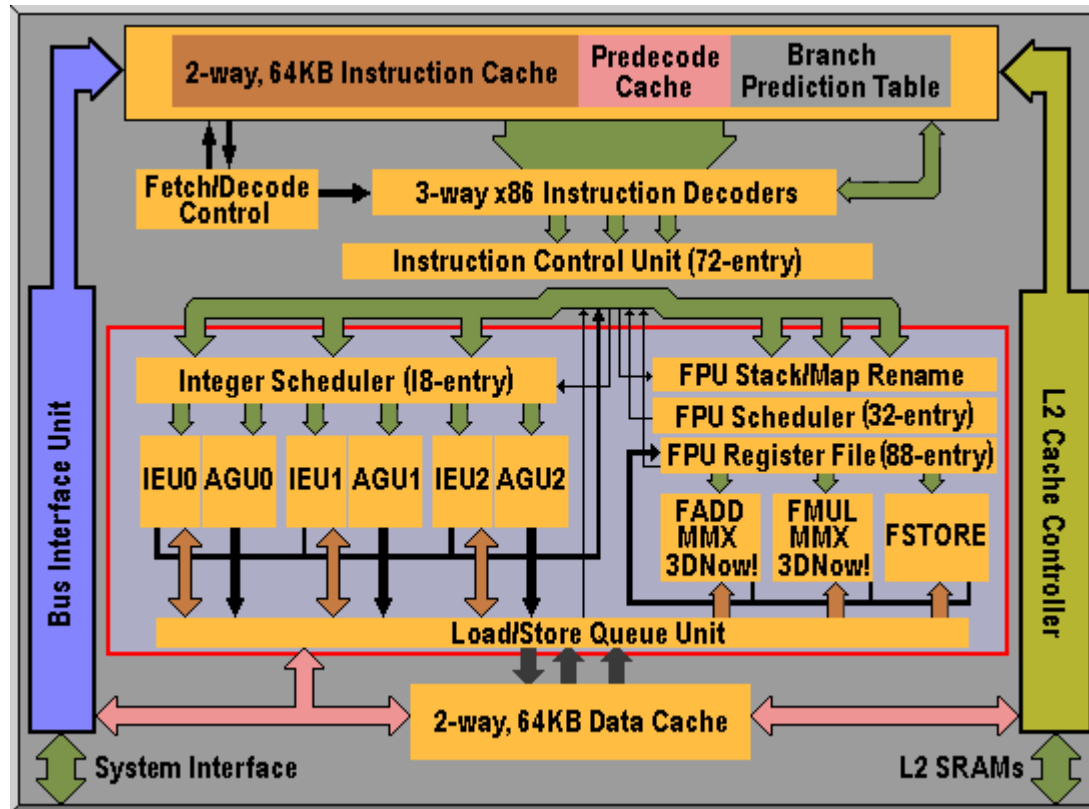


Pipelining



- Since there are different devices on the CPU that retrieve, decode, execute and store these can be run in parallel ("assembly line")
 - Can start retrieving a new instruction as soon as it is left for decoding.
- In modern Intel / AMD CPU there are «deep and long» pipelines.
 - 15-25 steps ->

Super scale




AMD Athlon hadde tre ALUer
og en SIMD flyttallenhet

- Superscalarity is a form of parallelization where we have **several** (parallel) **ALU** and **FPU** available on CPU
- Can then perform (even) more instructions at the same time!

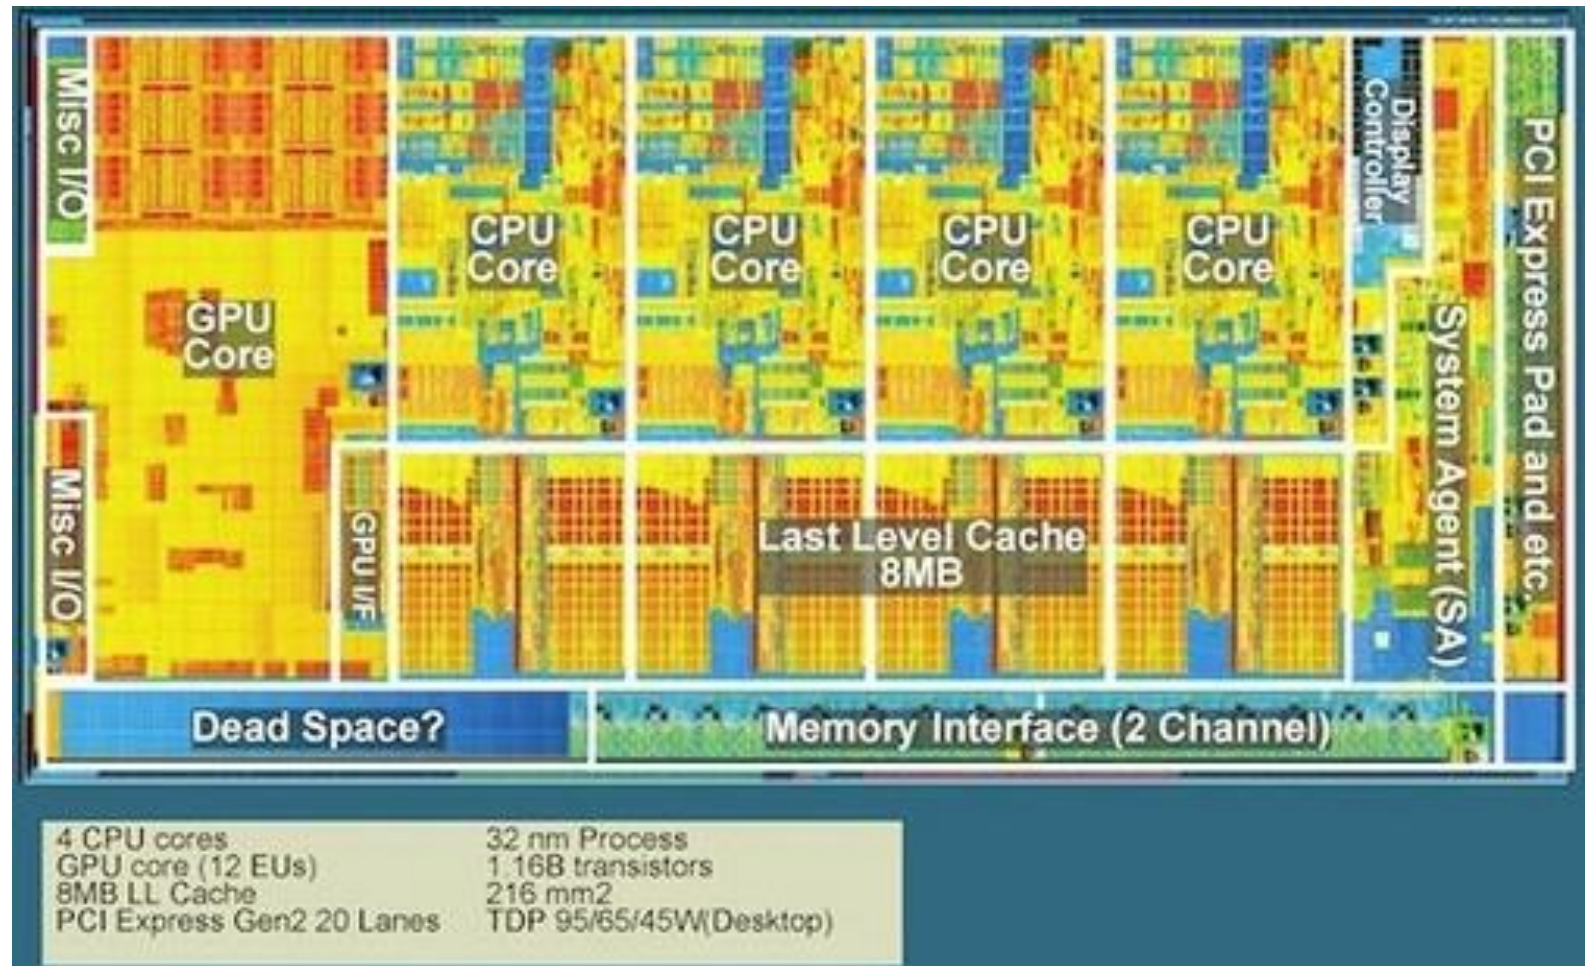
Pentium (IA32) layout (simplified)

- Branch Prediction
 - «"Guess" which instructions to run afterwards and retrieve them on the CPU»
- Cacheing
 - **Instructions** in own cache (should never return to RAM!)
 - **Data** in own cache

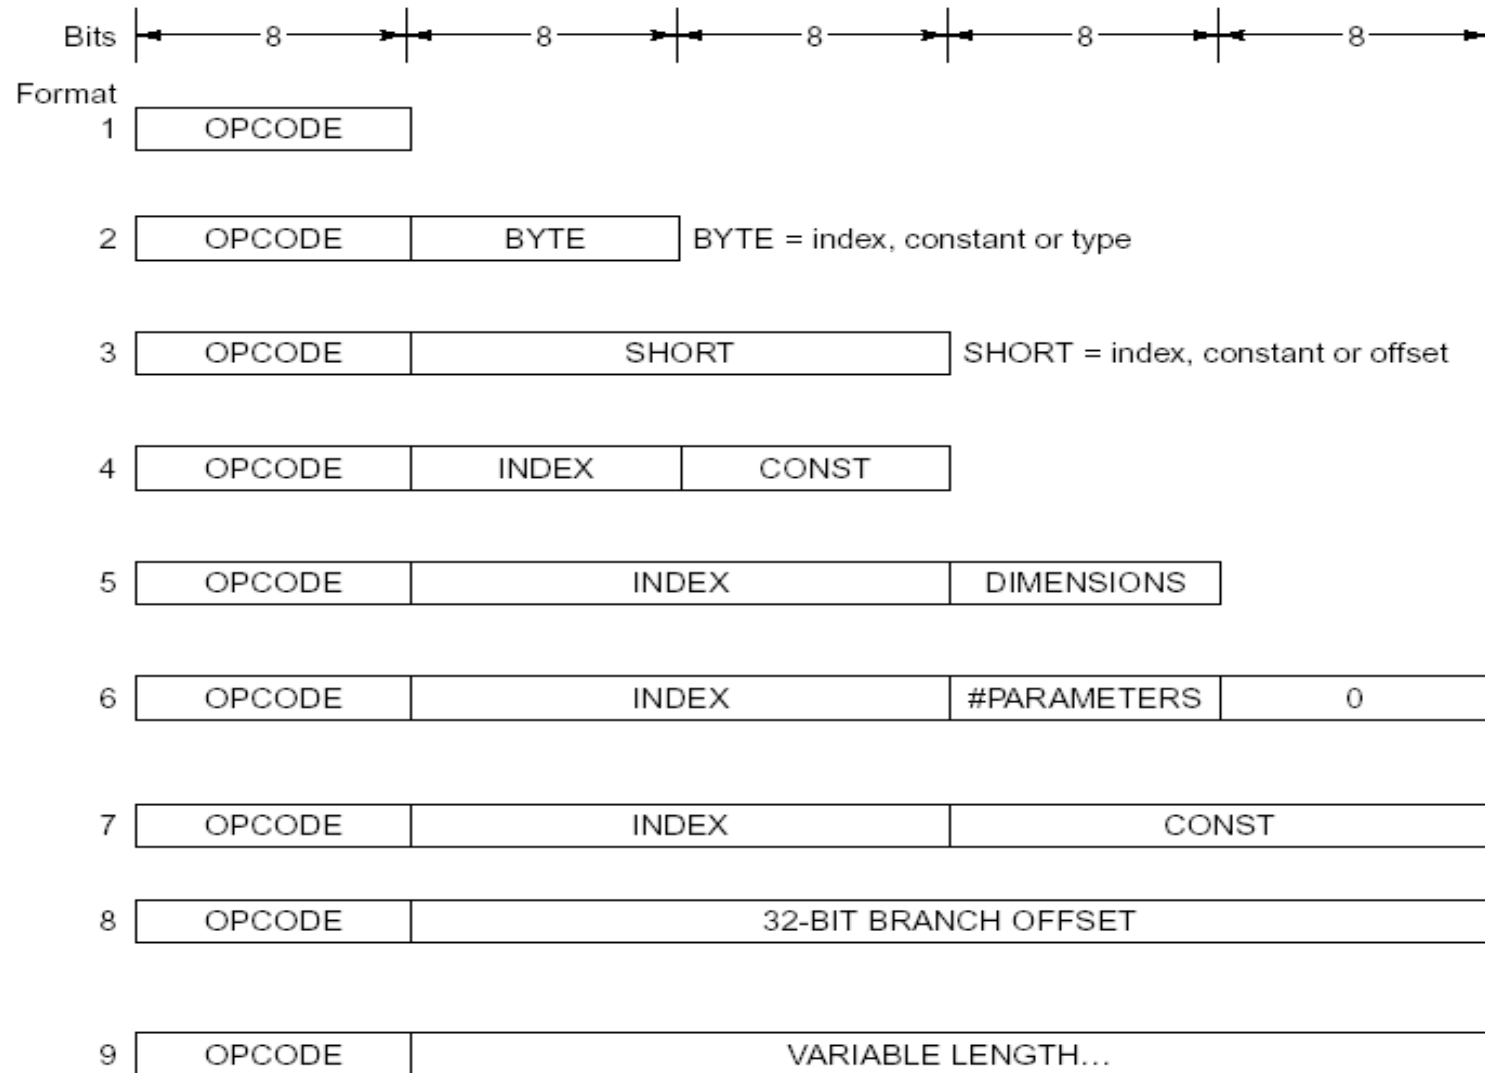


Sandy Bridge

- Multiple "cores" -> parallel running of program (threads)
- Moved Northbridge (and GPU) «on die»



JVM ISA: Format

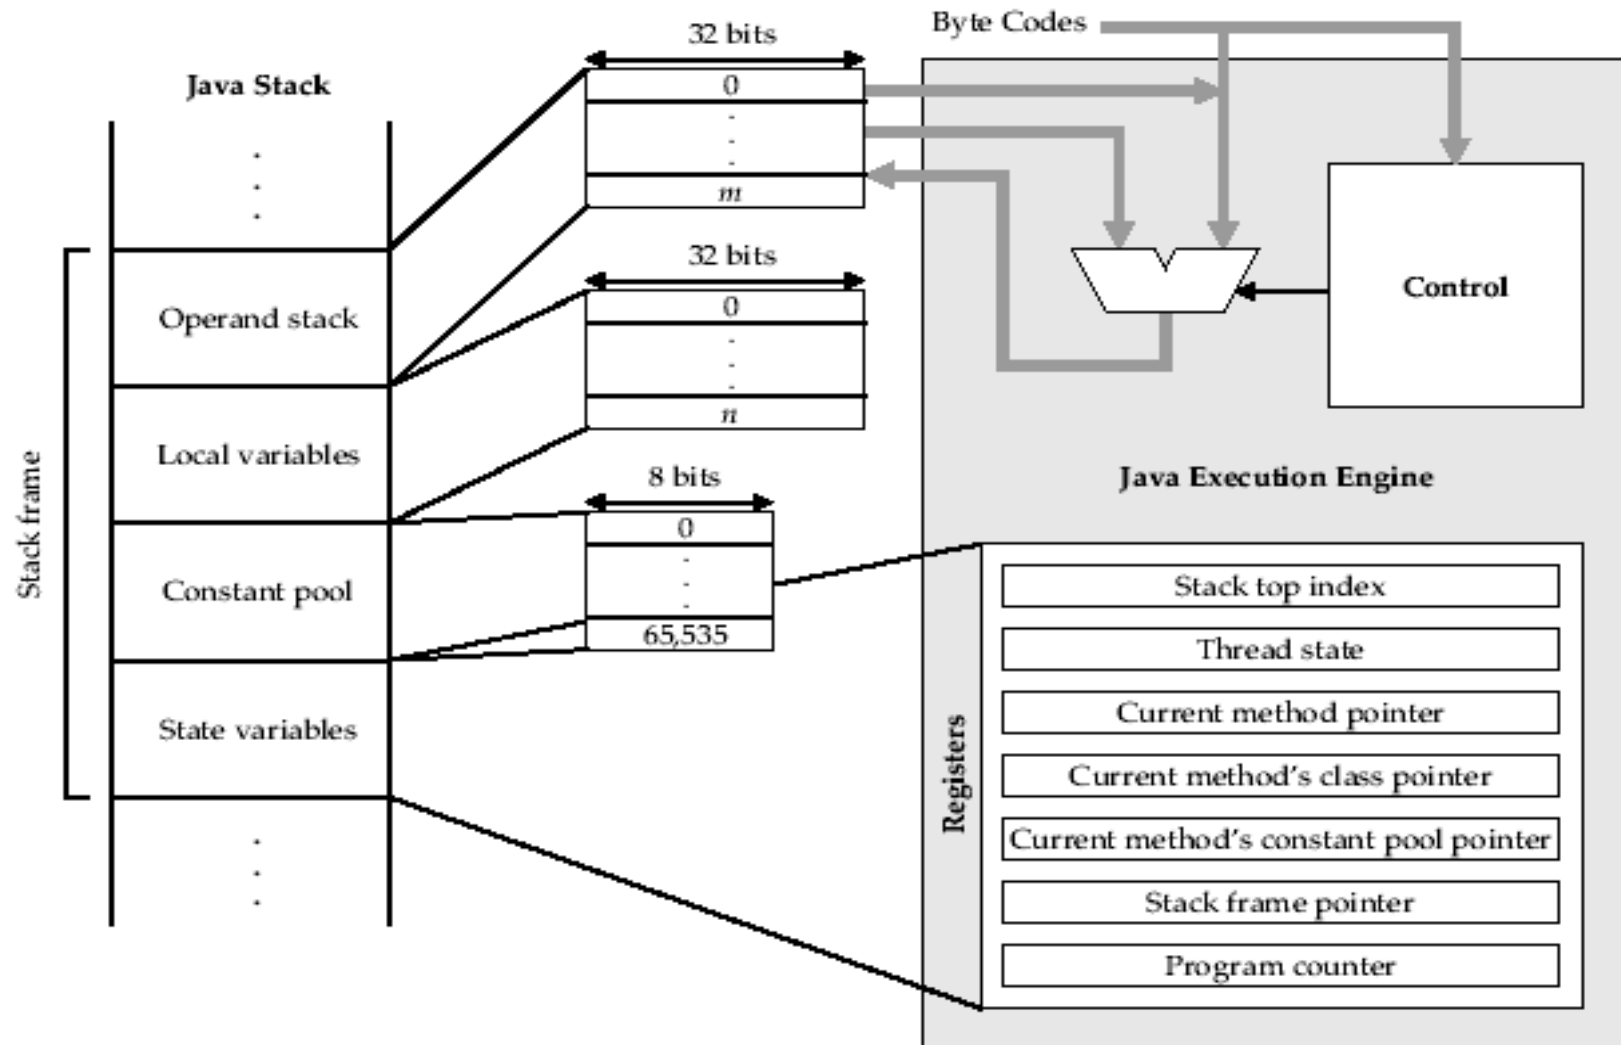


JVM: 226 instructions

Loads		Boolean/shift	
typeLOAD IND8	Push local variable onto stack	iIAND	Boolean AND
typeALOAD	Push array element on stack	iIOR	Boolean OR
BALOAD	Push byte from an array on stack	iIXOR	Boolean EXCLUSIVE OR
SALOAD	Push short from an array on stack	iISHL	Shift left
CALOAD	Push char from an array on stack	iISHR	Shift right
AALOAD	Push pointer from an array on "	iIUSHR	Unsigned shift right
Stores		Conversion	
typeSTORE IND8	Pop value and store in local var	x2y	Convert x to y
typeASTORE	Pop value and store in array	i2c	Convert integer to char
BASTORE	Pop byte and store in array	i2b	Convert integer to byte
SASTORE	Pop short and store in array	Stack management	
CASTORE	Pop char and store in array	DUPxx	Six instructions for duping
AASTORE	Pop pointer and store in array	POP	Pop an int from stk and discard
Pushes		POP2	Pop two ints from stk and discard
BIPUSH CON8	Push a small constant on stack	SWAP	Swap top two ints on stack
SIPUSH CON16	Push 16-bit constant on stack	Miscellaneous	
LDC IND8	Push constant from const pool	IINC IND8,CON8	Increment local variable
typeCONST_#	Push immediate constant	WIDE	Wide prefix
ACONST_NULL	Push a null pointer on stack	NOP	No operation
Arithmetic		GETFIELD IND16	Read field from object
typeADD	Add	PUTFIELD IND16	Write field to object
typeSUB	Subtract	GETSTATIC IND16	Get static field from class
typeMUL	Multiple	NEW IND16	Create a new object
typeDIV	Divide	INSTANCEOF OFFSET16	Determine type of obj
typeREM	Remainder	CHECKCAST IND16	Check object type
typeNEG	Negate	ATHROW	Throw exception
		LOOKUPSWITCH ...	Sparse multiway branch
		TABLESWITCH ...	Dense multiway branch
		MONITORENTER	Enter a monitor
		MONITOREXIT	Leave a monitor
Comparison		Transfer of control	
IF_ICMPrel OFFSET16	Conditional branch	INVOKEVIRTUAL IND16	Method invocation
IF_ACMPEQ OFFSET16	Branch if two ptrs equal	INVOKESTATIC IND16	Method invocation
IF_ACMPPNE OFFSET16	Branch if ptrs unequal	INVOKEINTERFACE ...	Method invocation
IFrel OFFSET16	Test 1 value and branch	INVOKESPECIAL IND16	Method invocation
IFNULL OFFSET16	Branch if ptr is null	JSR OFFSET16	Invoke finally clause
IFNONNULL OFFSET16	Branch if ptr is nonnull	typeRETURN	Return value
LCMP	Compare two longs	ARETURN	Return pointer
FCMPL	Compare 2 floats for <	RETURN	Return void
FCMPG	Compare 2 floats for >	RET IND8	Return from finally
DCMPL	Compare doubles for <	GOTO OFFSET16	Unconditional branch
DCMPG	Compare doubles for >	Arrays	
		ANEWARRAY IND16	Create array of ptrs
		NEWARRAY ATYPE	Create array of atype
		MULTINEWARRAY IN16,D	Create multidim array
		ARRAYLENGTH	Get array length

IND8/16 = index of local variable type, x, y = I, L, F, D
CON8/16, D, ATYPE = constant OFFSET16 for branch

Example: JVM (Java RunTime Enviroment)

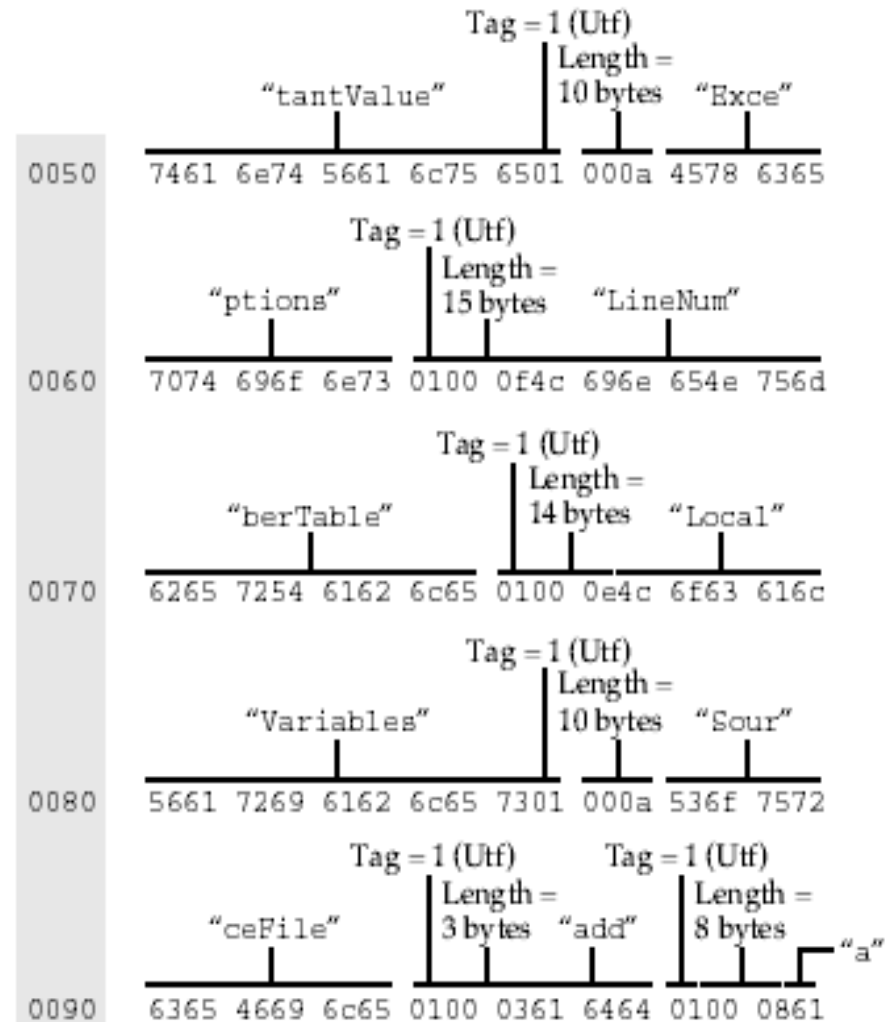
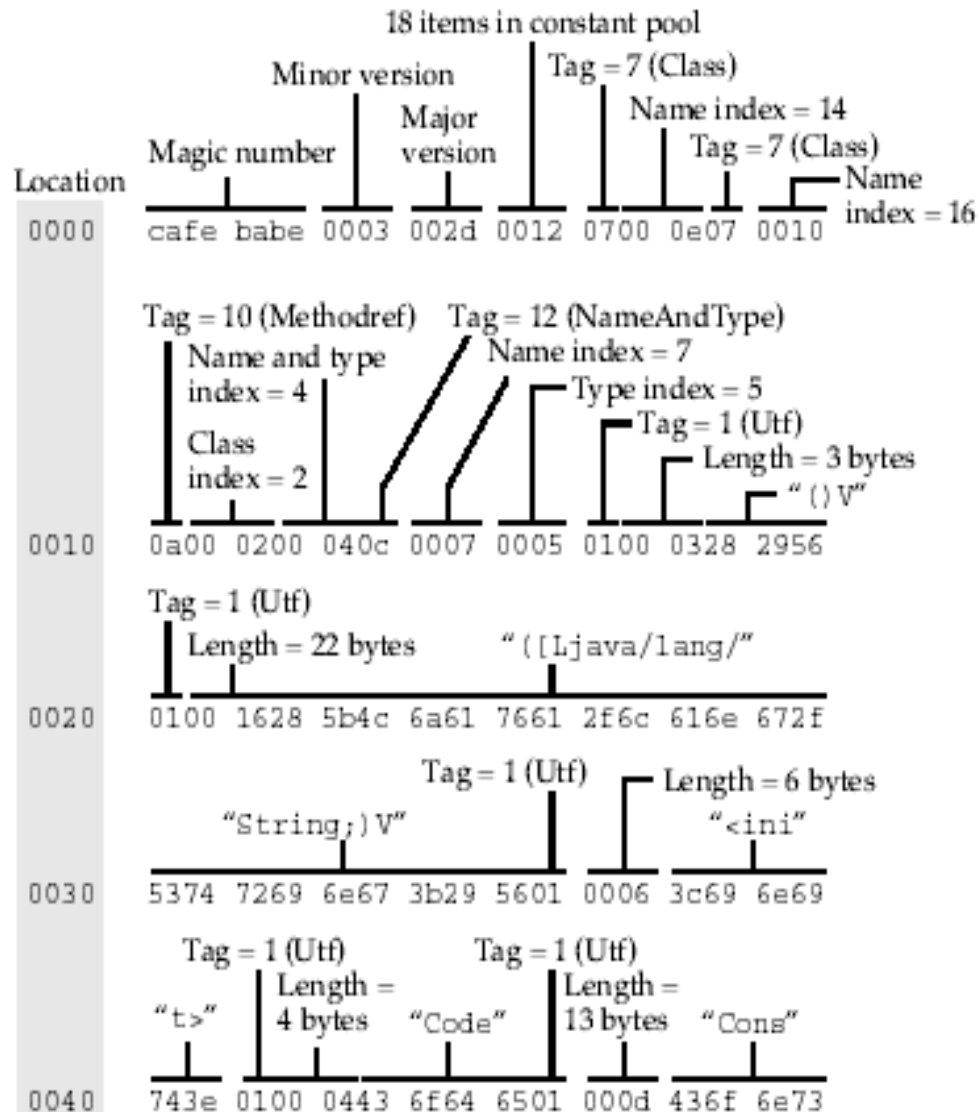


Example: Code and class

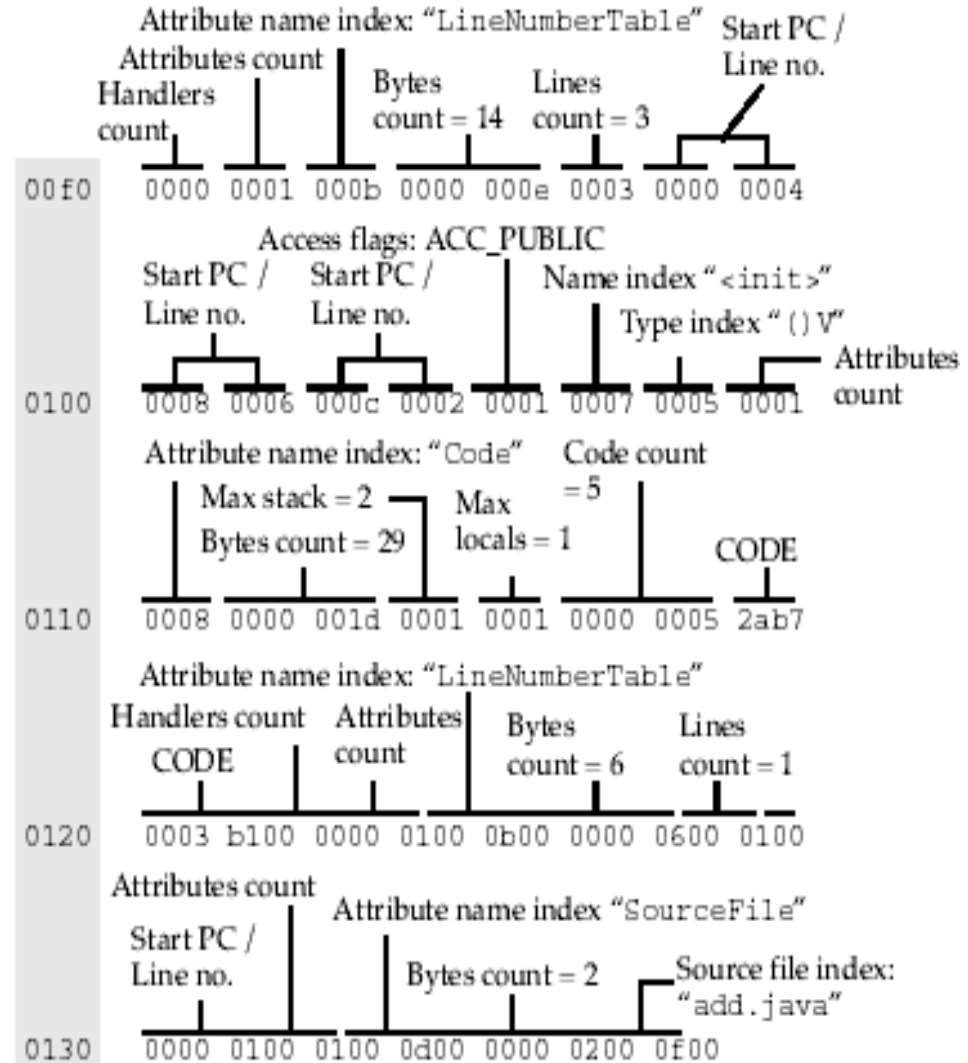
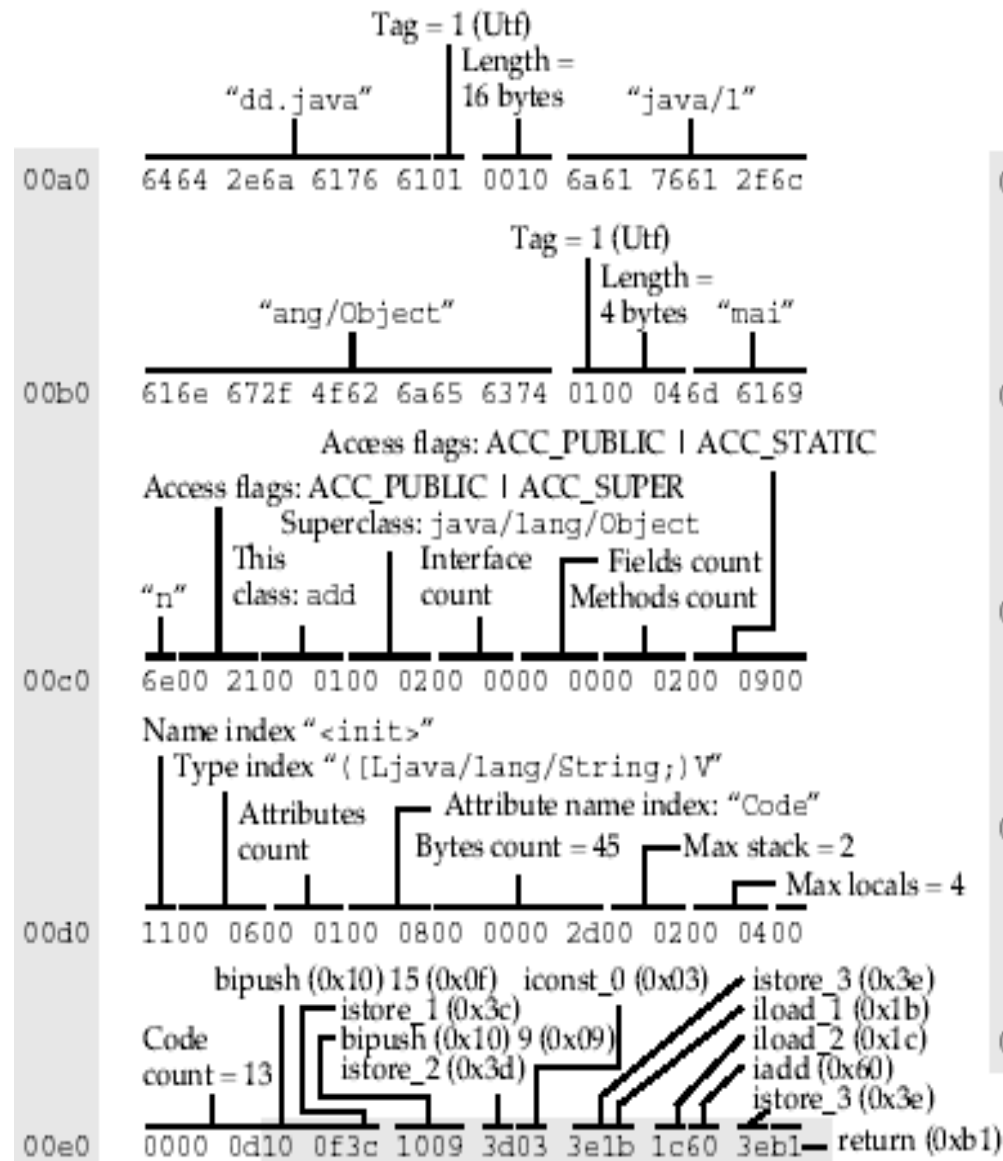
```
public class add {  
    public static void main(String args[]) {  
        int x=15, y=9, z=0;  
        z = x + y;  
    }  
}
```

0000	cafe	babe	0003	002d	0012	0700	0e07	0010
0010	0a00	0200	040c	0007	0005	0100	0328	2956()V
0020	0100	1628	5b4c	6a61	7661	2f6c	616e	672f	...([Ljava/lang/
0030	5374	7269	6e67	3b29	5601	0006	3c69	6e69	String;)V...<ini
0040	743e	0100	0443	6f64	6501	000d	436f	6e73	t>...Code...Cons
0050	7461	6e74	5661	6c75	6501	000a	4578	6365	tantValue...Exce
0060	7074	696f	6e73	0100	0f4c	696e	654e	756d	ptions...LineNum
0070	6265	7254	6162	6c65	0100	0e4c	6f63	616c	berTable...Local
0080	5661	7269	6162	6c65	7301	000a	536f	7572	Variables...Sour
0090	6365	4669	6c65	0100	0361	6464	0100	0861	ceFile...add...a
00a0	6464	2e6a	6176	6101	0010	6a61	7661	2f6c	dd.java...java/l
00b0	616e	672f	4f62	6a65	6374	0100	046d	6169	ang/Object...mai
00c0	6e00	2100	0100	0200	0000	0000	0200	0900	n.....
00d0	1100	0600	0100	0800	0000	2d00	0200	0400
00e0	0000	0d10	0f3c	1009	3d03	3e1b	1c60	3eb1
00f0	0000	0001	000b	0000	000e	0003	0000	0004
0100	0008	0006	000c	0002	0001	0007	0005	0001
0110	0008	0000	001d	0001	0001	0000	0005	2ab7
0120	0003	b100	0000	0100	0b00	0000	0600	0100
0130	0000	0100	0100	0d00	0000	0200	0f00	

Class file (1)



Class file (2)



Byte code

<u>Location</u>	<u>Code</u>	<u>Mnemonic</u>	<u>Meaning</u>
0x00e3	0x10	bipush	Push next byte onto stack
0x00e4	0x0f	15	Argument to bipush
0x00e5	0x3c	istore_1	Pop stack to local variable 1
0x00e6	0x10	bipush	Push next byte onto stack
0x00e7	0x09	9	Argument to bipush
0x00e8	0x3d	istore_2	Pop stack to local variable 2
0x00e9	0x03	iconst_0	Push 0 onto stack
0x00ea	0x3e	istore_3	Pop stack to local variable 3
0x00eb	0x1b	iload_1	Push local variable 1 onto stack
0x00ec	0x1c	iload_2	Push local variable 2 onto stack
0x00ed	0x60	iadd	Add top two stack elements
0x00ee	0x3e	istore_3	Pop stack to local variable 3
0x00ef	0xb1	return	Return

Java byte code vs x86 instructions

- Java bytecode can be called a **RISC** type instruction set
- x86 (Intel, as well as AMD) is **CISC**
- The next two files show:
 - The oldest instructions still included (256 of them, there are many, many more more)
 - Example of an IA-32 instruction

X86 – Basic Encoding Room

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0	ADD rmb,rb	ADD rmv,rv	ADD rb,rmb	ADD rv,rmv	ADD AL,ib	ADD eAX,iv	PUSH ES	POP ES	OR rmb,rb	OR rmv,rv	OR rb,rmb	OR rv,rmv	OR AL,ib	OR eAX,iv	PUSH CS	386 space
1	ADC rmb,rb	ADC rmv,rv	ADC rb,rmb	ADC rv,rmv	ADC AL,ib	ADC eAX,iv	PUSH SS	POP SS	SBB rmb,rb	SBB rmv,rv	SBB rb,rmb	SBB rv,rmv	SBB AL,ib	SBB eAX,iv	PUSH DS	POP DS
2	AND rmb,rb	AND rmv,rv	AND rb,rmb	AND rv,rmv	AND AL,ib	AND eAX,iv	ES:	DAA	SUB rmb,rb	SUB rmv,rv	SUB rb,rmb	SUB rv,rmv	SUB AL,ib	SUB eAX,iv	CS:	DAS
3	XOR rmb,rb	XOR rmv,rv	XOR rb,rmb	XOR rv,rmv	XOR AL,ib	XOR eAX,iv	SS:	AAA	CMP rmb,rb	CMP rmv,rv	CMP rb,rmb	CMP rv,rmv	CMP AL,ib	CMP eAX,iv	DS:	AAS
4	INC eAX	INC eCX	INC eDX	INC eBX	INC eSP	INC eBP	INC eSI	INC eDI	DEC eAX	DEC eCX	DEC eDX	DEC eBX	DEC eSP	DEC eBP	DEC eSI	DEC eDI
5	PUSH eAX	PUSH eCX	PUSH eDX	PUSH eBX	PUSH eSP	PUSH eBP	PUSH eSI	PUSH eDI	POP eAX	POP eCX	POP eDX	POP eBX	POP eSP	POP eBP	POP eSI	POP eDI
6	PUSHA PUSHAD	POPA POPAD	BOUND rv,m2v	ARPL rm2,r2	FS:	GS:	OpLen	AdLen	PUSH iv	IMUL rv,rmv,iv	PUSH ib	IMUL rv,rmv,ib	INSB	INSD	OUTSB	OUTSD
7	JO ib	JNO ib	JB ib	JAE ib	JE ib	JNE ib	JBE ib	JA ib	JS ib	JNS ib	JP ib	JNP ib	JL ib	JGE ib	JLE ib	JG ib
8	Immed rmb,ib	Immed rmv,iv		Immed rmv,ib	TEST rmb,rb	TEST rmv,rv	XCHG rmb,rb	XCHG rmv,rv	MOV rmb,rb	MOV rmv,rv	MOV rb,rmb	MOV rv,rmv	MOV rm,segr	LEA rv,m	MOV segr,rm	POP rmv
9	NOP	XCHG eAX,eCX	XCHG eAX,eDX	XCHG eAX,eBX	XCHG eAX,eSP	XCHG eAX,eBP	XCHG eAX,eSI	XCHG eAX,eDI	CWDE CBW	CWQ CDQ	CALL FAR sm	FWAIT	PUSHF	POPF	SAHF	LAHF
A	MOV AL,[iv]	MOV eAX,[iv]	MOV [iv],AL	MOV [iv],eAX	MOVSB	MOVSD	CMPBSB	CMPBSD	TEST AL,rb	TEST eAX,iv	STOSB	STOSD	LODSB	LODSD	SCASB	SCASD
B	MOV AL,ib	MOV CL,ib	MOV DL,ib	MOV BL,ib	MOV AH,ib	MOV CH,ib	MOV DH,ib	MOV BH,ib	MOV eAX,iv	MOV eCX,iv	MOV eDX,iv	MOV eBX,iv	MOV eSP,iv	MOV eBP,iv	MOV eSI,iv	MOV eDI,iv
C	Shift rmb,ib	Shift rmv,ib	RET i2	RET	LES rm,rmp	LDS rm,rmp	MOV rmb,ib	MOV rmv,iv	ENTER i2,ib	LEAVE	RETF i2	RETF	INT 3	INT ib	INTO	IRET
D	Shift rmb,1	Shift rmv,1	Shift rmb,CL	Shift rmv,CL	AAM	AAD		XLATB	87 space	87 space	87 space	87 space	87 space	87 space	87 space	87 space
E	LOOPNE short	LOOPE short	LOOP short	JecxZ short	IN AL,[ib]	IN eAX,[ib]	OUT [ib],AL	OUT [ib],eAX	CALL iv	JMP iv	JMP FAR sm	JMP ib	IN AL,[DX]	IN eAX,[DX]	OUT [DX],AL	OUT [DX],eAX
F	LOCK		REPNE	REP REPE	HLT	CMC	Unary rmb	Unary rmv	CLC	STC	CLI	STI	CLD	STD	IncDec rmb	Indir rmv

MOV (386->) – mod-reg-r/m

d and **w** in encoding define the length and direction of the operation



note: displacement may be zero, one, or two bytes long.

Table 24: MOD Encoding

MOD	Meaning
00	The r/m field denotes a register indirect memory addressing mode or a base/indexed addressing mode (see the encodings for r/m) <i>unless</i> the r/m field contains 110. If MOD=00 and r/m=110 the mod and r/m fields denote displacement-only (direct) addressing.
01	The r/m field denotes an indexed or base/indexed/displacement addressing mode. There is an eight bit signed displacement following the mod/reg/rm byte.
10	The r/m field denotes an indexed or base/indexed/displacement addressing mode. There is a 16 bit signed displacement (in 16 bit mode) or a 32 bit signed displacement (in 32 bit mode) following the mod/reg/rm byte .
11	The r/m field denotes a register and uses the same encoding as the <i>reg</i> field

Table 23: REG Bit Encodings

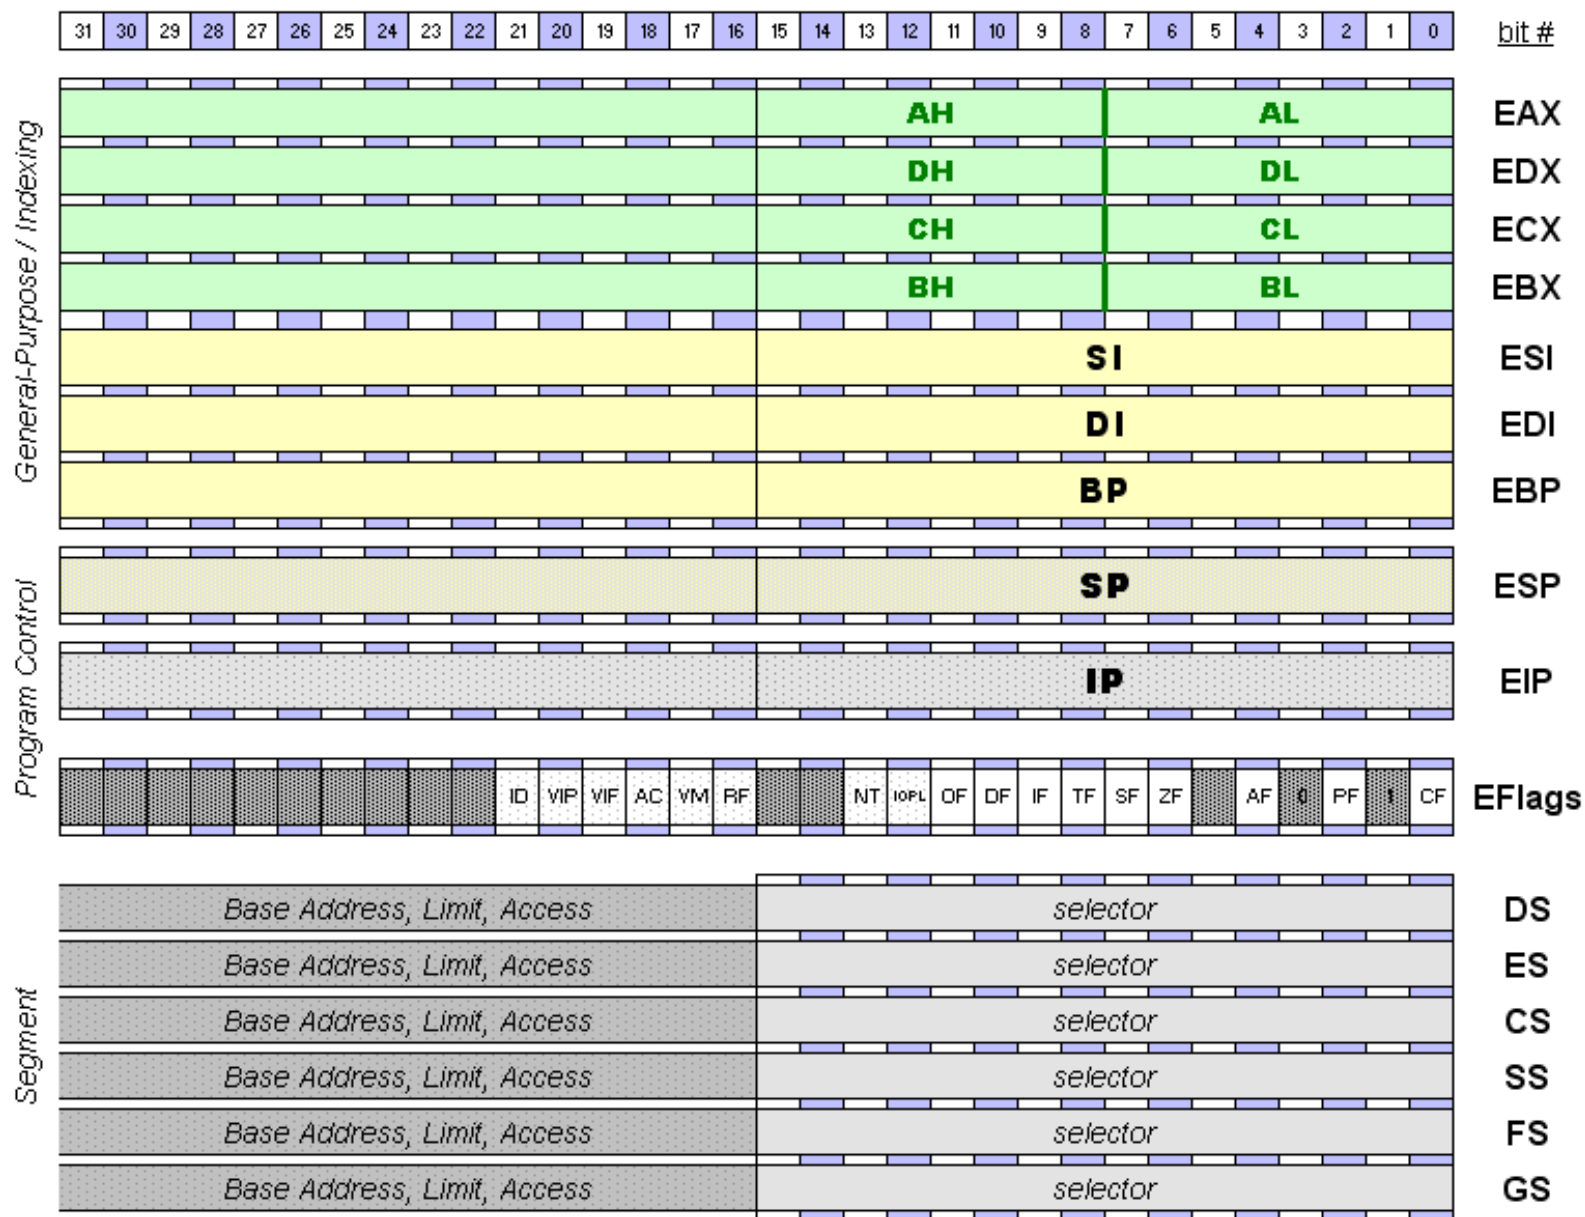
reg	w=0	16 bit mode w=1	32 bit mode w=1
000	AL	AX	EAX
001	CL	CX	ECX
010	DL	DX	EDX
011	BL	BX	EBX
100	AH	SP	ESP
101	CH	BP	EBP
110	DH	SI	ESI
111	BH	DI	EDI

Table 25: R/M Field Encoding


R/M	Addressing mode (Assuming MOD=00, 01, or 10)
000	[BX+SI] or DISP[BX][SI] (depends on MOD)
001	[BX+DI] or DISP[BX+DI] (depends on MOD)
010	[BP+SI] or DISP[BP+SI] (depends on MOD)
011	[BP+DI] or DISP[BP+DI] (depends on MOD)
100	[SI] or DISP[SI] (depends on MOD)
101	[DI] or DISP[DI] (depends on MOD)
110	Displacement-only or DISP[BP] (depends on MOD)
111	[BX] or DISP[BX] (depends on MOD)

[ebx] [ebp] ;Uses DS by default.
 [ebp] [ebx] ;Uses SS by default.
 [ebp*1] [ebx] ;Uses DS by default.
 [ebx] [ebp*1] ;Uses DS by default.
 [ebp] [ebx*1] ;Uses SS by default.
 [ebx*1] [ebp] ;Uses SS by default.
 es: [ebx] [ebp*1] ;Uses ES.

386 Registers (IA32) User Programs



X86-64 register



- Very many different registers -on each core!
- Extreme complexity
- How do we relate to complexity?
- **Abstract!**

- a concept or idea that is not associated with any specific instance
- to remove the unnecessary in a given context
- to focus on the features that are common
- To "hide" complexity

Example:

- In a Java program you can hide a complicated program inside a method...
- after that you just need to relate to the method
- High-level language abstracts away all the machine-complexity
- `System.out.print("Hello World!")` Actually involves a few thousand instructions since is a method call from `java.exe` to the operating system's write-to-screen routines.

What is abstraction?

- Abstraction is one of the four cornerstones of Computer Science. It involves filtering out – essentially, ignoring - the characteristics that we don't need in order to concentrate on those that we do.

